

Lecture Notes for **Machine Learning in Python**

Professor Eric Larson
Town Hall + MLP History

Class Logistics and Agenda

- Logistics:
 - Next time: Flipped Module on back propagation
- Multi Week Agenda:
 - Today: Neural Networks History, up to 1980
 - Today: Multi-layer Architectures
 - Town Hall, Lab 3 (if time)
 - Flipped: Programming Multi-layer training

A History of Neural Networks




Machine Learning 101


Neurons

- From biology to modeling:


input from neighboring neurons



McCulloch and Pitts, 1943



dendrite



input

logic gates of the mind




Warren McCulloch



Walter Pitts

Neurons

- McCulloch and Pitts, 1943
- Donald Hebb, 1949
 - Hebb's Law: close neurons fire together
 - neurons "learn"
 - easier synaptic
 - basis of neural




I as infatuated with the idea of **brainwashing** and controlling minds of others! I also invented a number of **torture procedures** like sensory deprivation and **isolation tanks**—and carried out a number of secret studies on real people!!



Donald O. Hebb



Warren McCulloch




Walter Pitts


Rosenblatt's perceptron, 1957



Frank Rosenblatt




dendrite



input

axon



hard limit



$$\begin{aligned} a &= -1 & z < 0 \\ a &= 1 & z \geq 0 \end{aligned}$$

linear



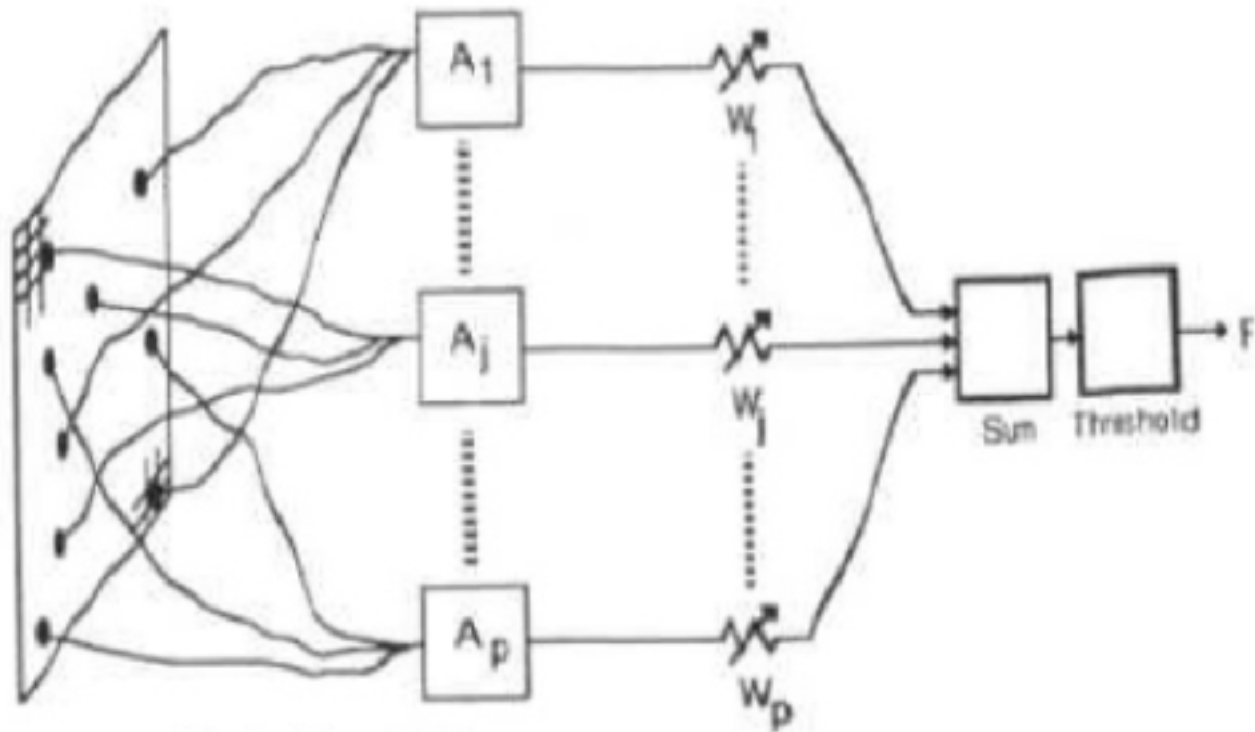
$$a = z$$

sigmoid

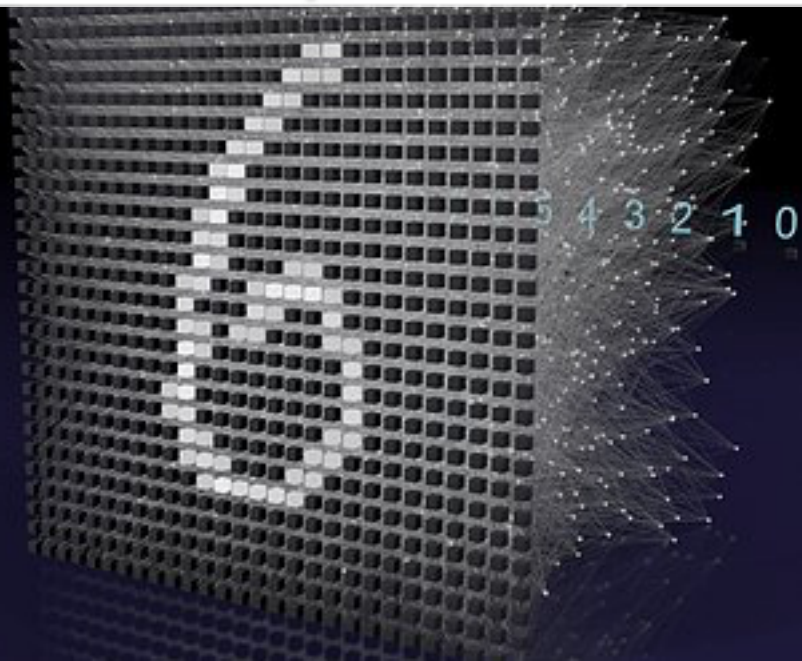


$$a = \frac{1}{1 + \exp(-z)}$$

The Mark 1




PERCEPTRON



Perceptron Learning Rule:
~Stochastic Gradient Descent

Layers Notation

dendrite



$\mathbf{a} = \mathbf{x} + \text{concat bias term}$

$$\mathbf{x}^{(i)} = \begin{bmatrix} x_1 \\ \vdots \\ x_j \\ \vdots \\ x_N \end{bmatrix}^{(i)} \quad \mathbf{a} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_j \\ \vdots \\ x_N \end{bmatrix}$$


$$\mathbf{z} = \mathbf{W} \cdot \mathbf{a} = \mathbf{W}_{1:N} \cdot \mathbf{x}^{(i)} + \mathbf{b}$$

$$\mathbf{W} = \begin{bmatrix} W_{1,0} & W_{1,1} & W_{1,2} & \dots & W_{1,N} \\ \dots & \dots & \dots & \dots & \dots \\ W_{S,0} & W_{S,1} & W_{S,2} & \dots & W_{S,N} \end{bmatrix}$$

$$[\mathbf{z}^{(1)}]^{(i)} = \mathbf{W}^{(1)} \cdot [\mathbf{a}^{(1)}]^{(i)} = \mathbf{W}_{1:N}^{(1)} \cdot \mathbf{x}^{(i)} + \mathbf{b}^{(1)}$$

$\mathbf{a}^{(\text{next})} = \boldsymbol{\phi}(\mathbf{z}^{(\text{current})}) + \text{concat bias term}$


$$\mathbf{a}^{(\text{next})} = \begin{bmatrix} 1 \\ \phi(z_1^{\text{curr}}) \\ \vdots \\ \phi(z_N^{\text{curr}}) \end{bmatrix} \rightarrow \mathbf{a}^{(L)} = \begin{bmatrix} 1 \\ \phi(z_1^{L-1}) \\ \vdots \\ \phi(z_N^{L-1}) \end{bmatrix}$$



$\mathbf{x}^{(i)}$ One row from Table data
becomes input column to model

notation adapted from *Neural Network Design*, Hagan, Demuth, Beale, and De Jesus

Multiple Layers Notation



$$\mathbf{a}^{(L+1)} = \phi(\mathbf{z}^{(L)}) + \text{concat bias term}$$

$$\mathbf{a}^{(4)} \text{ rows} = \text{unique classes}$$

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \mathbf{a}^{(L)}$$


$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \phi(\mathbf{z}^{(L-1)})$$

$$\mathbf{W}^{(L)} = \begin{bmatrix} w^{(L)}_{1,0} & w^{(L)}_{1,1} & w^{(L)}_{1,2} & \dots & w^{(L)}_{1,S^{L-1}} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w^{(L)}_{S^L,0} & w^{(L)}_{S^L,1} & \dots & \dots & w^{(L)}_{S^L,S^{L-1}} \end{bmatrix}$$

$S^L \times (S^{L-1}+1)$

$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \phi(\mathbf{W}^{(L-1)} \mathbf{a}^{(L-1)})$$


Multiple layers notation



- **Self test:** How many parameters need to be trained in the above network?
 - A. $[(N+1) \times S^1] + [(S^1 + 1) \times S^2] + [(S^2 + 1) \times S^3]$
 - B. $|\mathbf{W}^{(1)}| + |\mathbf{W}^{(2)}| + |\mathbf{W}^{(3)}|$
 - C. can't determine from diagram
 - D. it depends on the sizes of intermediate variables, $\mathbf{z}^{(i)}$

notation adapted from *Neural Network Design*, Hagan, Demuth, Beale, and De Jesus 11

Compact feedforward notation



$$\mathbf{z}^{(L)} = \mathbf{W}^{(L)} \mathbf{a}^{(L)}$$

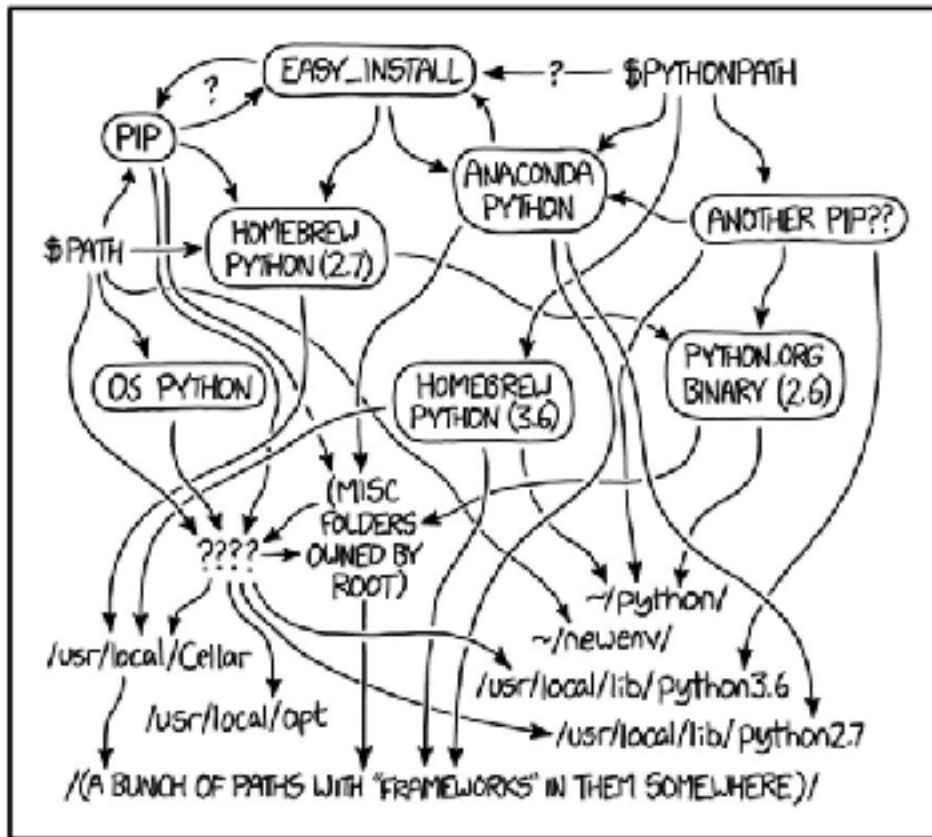
$$[\mathbf{z}^{(L)}]^{(i)} = \mathbf{W}^{(L)} [\mathbf{a}^{(L)}]^{(i)} \quad \begin{bmatrix} z^{(L)}_1 \\ \vdots \\ z^{(L)}_{S^L} \end{bmatrix}^{(i)} = \begin{bmatrix} w^{(L)}_{1,1} & w^{(L)}_{1,2} & w^{(L)}_{1,3} & \dots & w^{(L)}_{1,S^{L-1}+1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w^{(L)}_{S^L,1} & w^{(L)}_{S^L,2} & \dots & \dots & w^{(L)}_{S^L,S^{L-1}+1} \end{bmatrix} \begin{bmatrix} a^{(L)}_1 \\ \vdots \\ a^{(L)}_{S^{L-1}+1} \end{bmatrix}^{(i)}$$

$$\begin{bmatrix} \begin{bmatrix} z^{(L)}_1 \\ \vdots \\ z^{(L)}_{S^L} \end{bmatrix}^{(1)} & \dots & \begin{bmatrix} z^{(L)}_1 \\ \vdots \\ z^{(L)}_{S^L} \end{bmatrix}^{(M)} \end{bmatrix} = [\mathbf{W}^{(L)}] \begin{bmatrix} \begin{bmatrix} a^{(L)}_1 \\ \vdots \\ a^{(L)}_{S^{L-1}+1} \end{bmatrix}^{(1)} & \dots & \begin{bmatrix} a^{(L)}_1 \\ \vdots \\ a^{(L)}_{S^{L-1}+1} \end{bmatrix}^{(M)} \end{bmatrix}$$


$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \mathbf{A}^{(L)}$$

$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \phi(\mathbf{W}^{(L-1)} \mathbf{A}^{(L-1)})$$

Training Neural Network Architectures



Rosenblatt's Perceptron, 1957



$$\sum_i^M (\mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)})^2$$



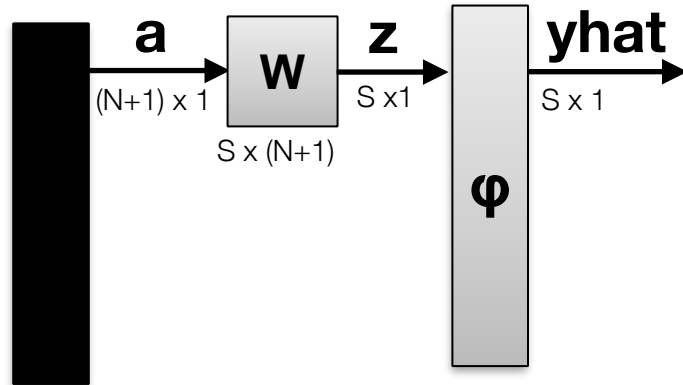
Need objective Function, minimize MSE $J(\mathbf{W}) = \|\mathbf{Y} - \hat{\mathbf{Y}}\|^2$

where ground truth $\mathbf{y}^{(i)}$ is one-hot encoded!

$$\begin{bmatrix} y_1 \\ \dots \\ y_c \end{bmatrix}^{(i)} \rightarrow \begin{bmatrix} \begin{bmatrix} y_1 \\ \dots \\ y_c \end{bmatrix}^{(1)} \quad \dots \quad \begin{bmatrix} y_1 \\ \dots \\ y_c \end{bmatrix}^{(M)} \end{bmatrix} = \mathbf{Y}$$

Simple Architectures

- Rosenblatt's perceptron, 1957




Self Test - If this is a binary classification problem, how large is S , the length of $\mathbf{\hat{y}}$?


- A. Can't determine
- B. 2
- C. 1
- D. N

Simple Architectures

- Adaline network, Widrow and Hoff, 1960



Marcian "Ted" Hoff



Bernard Widrow

Objective Function, minimize MSE $J(\mathbf{W}) = ||\mathbf{Y} - \hat{\mathbf{Y}}||^2$



New objective function becomes: $J(\mathbf{W}) = ||\mathbf{Y} - \mathbf{W} \cdot \mathbf{A}||^2$

Need gradient $\nabla J(\mathbf{W})$ for update equation $\mathbf{W} \leftarrow \mathbf{W} + \eta \nabla J(\mathbf{W})$

We have been using the **Widrow-Hoff Learning Rule**

Simple Architectures

- Adaline network, Widrow and Hoff, 1960



Marcian "Ted" Hoff



Bernard Widrow

need gradient $\nabla J(\mathbf{W})$ for update equation $\mathbf{W} \leftarrow \mathbf{W} + \eta \nabla J(\mathbf{W})$


For case $S=1$, \mathbf{W} has only one row, \mathbf{w} this is just **linear regression**...

$$\mathbf{w} \leftarrow \mathbf{w} + \eta [\mathbf{X} * (\mathbf{y} - \hat{\mathbf{y}})]$$



Simple Architectures

- Modern Perceptron network



$$\phi(z) = \frac{1}{1 + \exp(-z)}$$

need gradient $\nabla J(\mathbf{W})$ for update equation $\mathbf{W} \leftarrow \mathbf{W} + \eta \nabla J(\mathbf{W})$


For case $S=1$, this is just **logistic regression...**
and **we have already solved this!**

$$\mathbf{w} \leftarrow \mathbf{w} + \eta [\mathbf{X} * (\mathbf{y} - \mathbf{g}(\mathbf{x}))]$$



What happens when $S > 1$?

What if we have more than $S=1$?



$$\begin{bmatrix} \begin{bmatrix} \phi(z_1) \\ \dots \\ \phi(z_s) \end{bmatrix}^{(1)} & \dots & \begin{bmatrix} \phi(z_1) \\ \dots \\ \phi(z_s) \end{bmatrix}^{(M)} \end{bmatrix} = \hat{\mathbf{Y}}^{\Delta}$$

$$\begin{bmatrix} \begin{bmatrix} y_1 \\ \dots \\ y_s \end{bmatrix}^{(1)} & \dots & \begin{bmatrix} y_1 \\ \dots \\ y_s \end{bmatrix}^{(M)} \end{bmatrix} = \mathbf{Y}$$

$$J(\mathbf{W}) = ||\mathbf{Y} - \hat{\mathbf{Y}}^{\Delta}||^2$$

Each target class in \mathbf{Y} can be independently optimized

$$\mathbf{yhat}^{(i)} = \begin{bmatrix} \phi(\text{row}=1 \mathbf{W} \cdot \mathbf{x}^{(i)}) \\ \phi(\text{row}=2 \mathbf{W} \cdot \mathbf{x}^{(i)}) \\ \dots \\ \phi(\text{row}=S \mathbf{W} \cdot \mathbf{x}^{(i)}) \end{bmatrix}$$

one hot



which is one-versus-all!

$$J(1\mathbf{W}) = \sum_{i=1} [y_1^{(i)} - \phi(1\mathbf{W} \cdot \mathbf{x}^{(i)})]^2$$

$$J(2\mathbf{W}) = \sum_{i=1} [y_2^{(i)} - \phi(2\mathbf{W} \cdot \mathbf{x}^{(i)})]^2$$

...

$$J(s\mathbf{W}) = \sum_{i=1} [y_s^{(i)} - \phi(s\mathbf{W} \cdot \mathbf{x}^{(i)})]^2$$

Simple Architectures: Summary

- Adaline network, Widrow and Hoff, 1960
 - linear regression
- Perceptron
 - *with sigmoid*: logistic regression
- One-versus-all implementation is the same as having $\mathbf{w}_{\text{class}}$ be rows of weight matrix, \mathbf{W}
 - works in adaline
 - works in logistic regression



these networks were created in the 50's and 60's
but were abandoned

why were they not used?

The Rosenblatt-Widrow-Hoff Dilemma

- 1960's: Rosenblatt got into a public academic argument with Marvin Minsky and Seymour Papert

"Given an elementary α -perceptron, a stimulus world W , and any classification $C(W)$ for which a solution exists; let all stimuli in W occur in any sequence, provided that each stimulus must reoccur in finite time; then beginning from an arbitrary initial state, an error correction procedure will always yield a solution to $C(W)$ in finite time..."

- Minsky and Papert publish limitations paper, 1969:

TED Ideas worth spreading

WATCH

DISCOVER

ATTEND

PARTICIPATE

Marvin Minsky:

Health and the human mind

TED2008 · 13:33 · Filmed Feb 2008

21 subtitle languages

View interactive transcript



More Advanced Architectures: history

- 1986: *Rumelhart, Hinton, and Williams* popularize gradient calculation for multi-layer network
 - *technically* introduced by Werbos in 1982
- **difference:** Rumelhart *et al.* validated ideas with a computer
- until this point no one could train a multiple layer network consistently
- algorithm is popularly called **Back-Propagation**
- wins pattern recognition prize in 1993, becomes de-facto machine learning algorithm until: SVMs and Random Forests in ~2004
- would eventually see a resurgence for its ability to train algorithms for Deep Learning applications: **Hinton is widely considered the founder of deep learning**

David Rumelhart

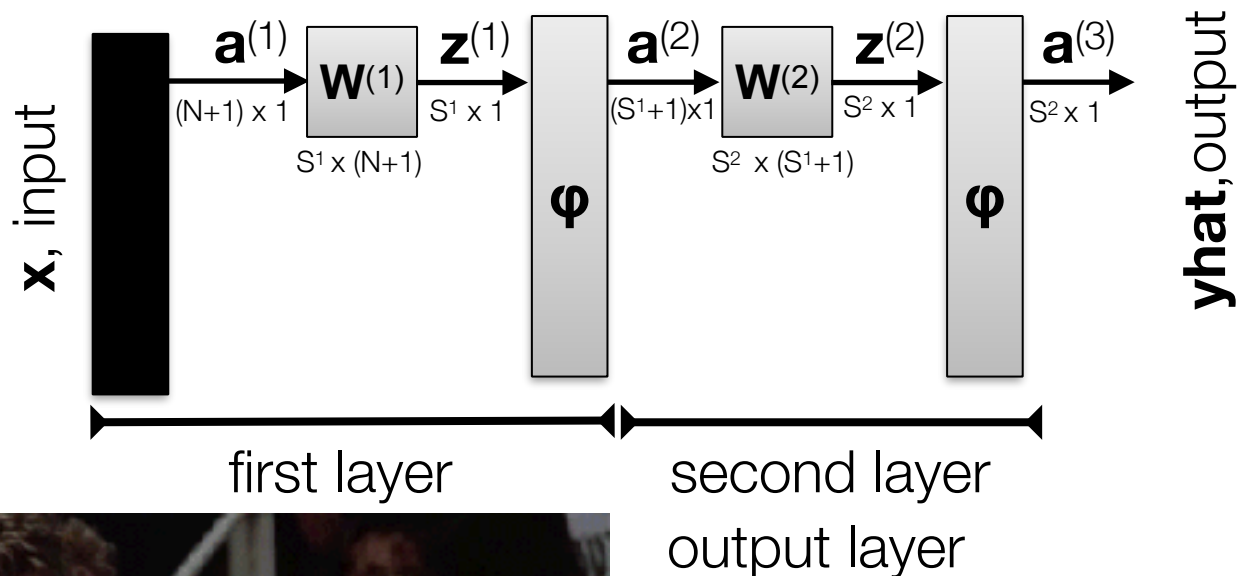


Geoffrey Hinton



More Advanced Architectures: MLP

- The multi-layer perceptron (MLP):
 - two layers shown, but could be arbitrarily many layers
 - algorithm is agnostic to number of layers (*kinda*)



each row of **yhat** is no longer independent of the rows in **W** so we cannot optimize using one versus all!!!




$$\mathbf{yhat}^{(i)} = \begin{bmatrix} \phi(\text{row}=1 \mathbf{W}^{(2)} \cdot \phi(\mathbf{W}^{(1)} \mathbf{a}^{(1)})) \\ \vdots \\ \phi(\text{row}=S \mathbf{W}^{(2)} \cdot \phi(\mathbf{W}^{(1)} \mathbf{a}^{(1)})) \end{bmatrix}$$

one hot

Back propagation

- Steps:
 - propagate weights forward
 - calculate gradient at final layer
 - back propagate gradient for each layer
 - via recurrence relation




$\hat{\mathbf{y}}$, output

$$J(\mathbf{W}) = \|\mathbf{Y} - \hat{\mathbf{Y}}\|^2$$

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta \frac{\partial J(\mathbf{W})}{\partial w_{ij}^{(l)}}$$

Back propagation



use chain rule:

$$\frac{\partial J(\mathbf{W})}{\partial w_{ij}^{(l)}} = \frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(l)}} \frac{\partial \mathbf{z}^{(l)}}{\partial w_{ij}^{(l)}}$$

$$J(\mathbf{W}) = \sum_k^M (\mathbf{y}^{(k)} - \mathbf{a}^{(L)})^2$$

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta \frac{\partial J(\mathbf{W})}{\partial w_{ij}^{(l)}}$$

Solve this in explainer video for next in class assignment!

Lab 3, Town Hall (if time)



End of Session

- thanks! **Next time is Flipped Assignment!!!**

More help on neural networks to prepare for next time:

Sebastian Raschka

<https://github.com/rasbt/python-machine-learning-book/blob/master/code/ch12/ch12.ipynb>

Martin Hagan

https://www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=1&cad=rja&uact=8&ved=0ahUKEwioprwn27fPAhWMx4MKHYbwDIwQFggeMAA&url=http%3A%2F%2Fhagan.okstate.edu%2FNNDesign.pdf&usg=AFQjCNG5YbM4xSMm6K5HNsG-4Q8TvOu_Lw&sig2=bgT3k-5ZDDTPZ07Qu8Oreg

Michael Nielsen

<http://neuralnetworksanddeeplearning.com>