

# Lecture Notes for **Machine Learning in Python**

Professor Eric Larson  
**Optimizing Neural Networks**

# Class Logistics and Agenda

- Logistics
- Agenda:
  - ▣ Finish Town Hall
  - ▣ Practical Multi-layer Architectures
  - ▣ Programming Examples
- ▣ Next Time: More MLPs



**Tyler Rablin** @Mr\_Rablin · 2d  
You're not grading assignments.

You're collecting evidence to determine student progress and pointing them towards their next steps.

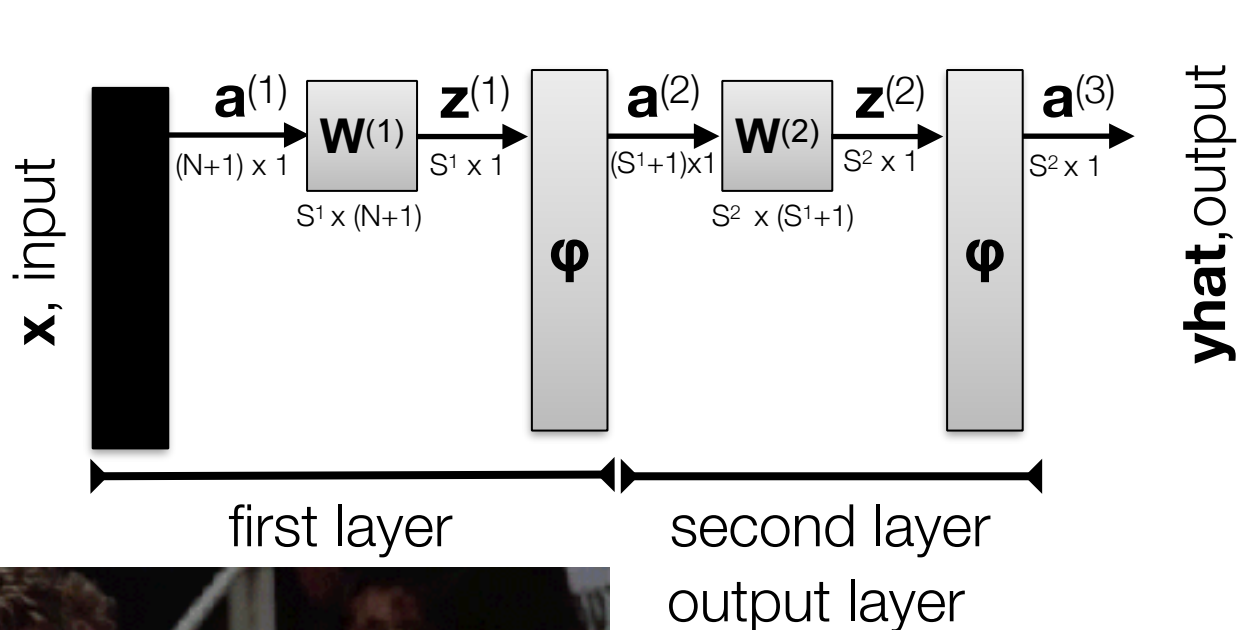
Make the mental switch. It matters.

# Town Hall



# More Advanced Architectures: MLP

- The multi-layer perceptron (MLP):
  - two layers shown, but could be arbitrarily many layers



each row of **yhat** is no longer independent of the rows in **W** so we cannot optimize using one versus all!!!



$$\mathbf{yhat}^{(i)} = \begin{bmatrix} \phi(\text{row}=1 \mathbf{W}^{(2)} \cdot \phi(\mathbf{W}^{(1)} \mathbf{a}^{(1)})) \\ \vdots \\ \phi(\text{row}=S \mathbf{W}^{(2)} \cdot \phi(\mathbf{W}^{(1)} \mathbf{a}^{(1)})) \end{bmatrix}$$

one hot


# The Rosenblatt-Widrow-Hoff Dilemma

- 1960's: Rosenblatt got into a public academic argument with Marvin Minsky and Seymour Papert

"Given an elementary  $\alpha$ -perceptron, a stimulus world  $W$ , and any classification  $C(W)$  for which a solution exists; let all stimuli in  $W$  occur in any sequence, provided that each stimulus must reoccur in finite time; then beginning from an arbitrary initial state, an error correction procedure will always yield a solution to  $C(W)$  in finite time..."

- Minsky and Papert publish limitations paper, 1969:

"the style of research being done on the perceptron is doomed to failure because of these limitations."



- Widrow and Rosenblatt try to build bigger networks without limitations and fail
  - ❓ Neural Networks research **basically stops** for **17 years**
- **Until:** researchers revisit training bigger networks
  - ❓ neural networks with multiple layers

# More Advanced Architectures: history

- 1986: *Rumelhart, Hinton, and Williams* popularize gradient calculation for multi-layer network
  - *actually* introduced by Werbos in 1982
- **difference:** Rumelhart *et al.* validated ideas with a computer
- until this point no one could train a multiple layer network consistently
- algorithm is popularly called **Back-Propagation**
- wins pattern recognition prize in 1993, becomes de-facto machine learning algorithm until: SVMs and Random Forests in ~2004
- would eventually see a resurgence for its ability to train algorithms for Deep Learning applications: **Hinton is widely considered the founder of deep learning**

David Rumelhart




Geoffrey Hinton



# Back propagation

- Steps:
  - propagate weights forward
  - calculate gradient at final layer
  - back propagate gradient for each layer
    - via recurrence relation




$$J(\mathbf{W}) = \|\mathbf{Y} - \hat{\mathbf{Y}}\|^2$$

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta \frac{\partial J(\mathbf{W})}{\partial w_{ij}^{(l)}}$$

\*\*Recall from Flipped Assignment!

# Back Propagation Summary



1. Forward propagate to get  $\mathbf{Z}$ ,  $\mathbf{A}$
2. Get final layer gradient
3. Back propagate sensitivities
4. Update each  $\mathbf{W}^{(l)}$

$$\mathbf{V}^{(2)} = -2(\mathbf{Y} - \mathbf{A}^{(3)}) * \mathbf{A}^{(3)} * (1 - \mathbf{A}^{(3)})$$
$$\nabla^{(2)} = \mathbf{V}^{(2)} \cdot [\mathbf{A}^{(2)}]^T$$

$$\mathbf{V}^{(1)} = \mathbf{A}^{(2)} * (1 - \mathbf{A}^{(2)}) * [\mathbf{W}^{(2)}]^T \cdot \mathbf{V}^{(2)}$$
$$\nabla^{(1)} = \mathbf{V}^{(1)} \cdot [\mathbf{A}^{(1)}]^T$$

$$\mathbf{W}^{(l)} \leftarrow \mathbf{W}^{(l)} - \eta \nabla^{(l)}$$


Where is the problem of  
**vanishing gradients** introduced?

\*\*Recall from Flipped Assignment!



## 07. MLP Neural Networks.ipynb

same as Flipped Assignment!  
with regularization  
and vectorization



# Problems with Advanced Architectures

- Numerous weights to find gradient update
  - minimize number of instances
  - **solution:** mini-batch
- **new problem:** mini-batch gradient can be erratic
  - **solution:** momentum
    - use previous update in current update

# Common Adaptive Strategies


$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Momentum

$$\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$$


- Nesterov's Accelerated Gradient

$$\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1}))}_{\text{step twice}} + \alpha \nabla J(\mathbf{W}_{k-1})$$



# Adaptive Strategy: Cooling

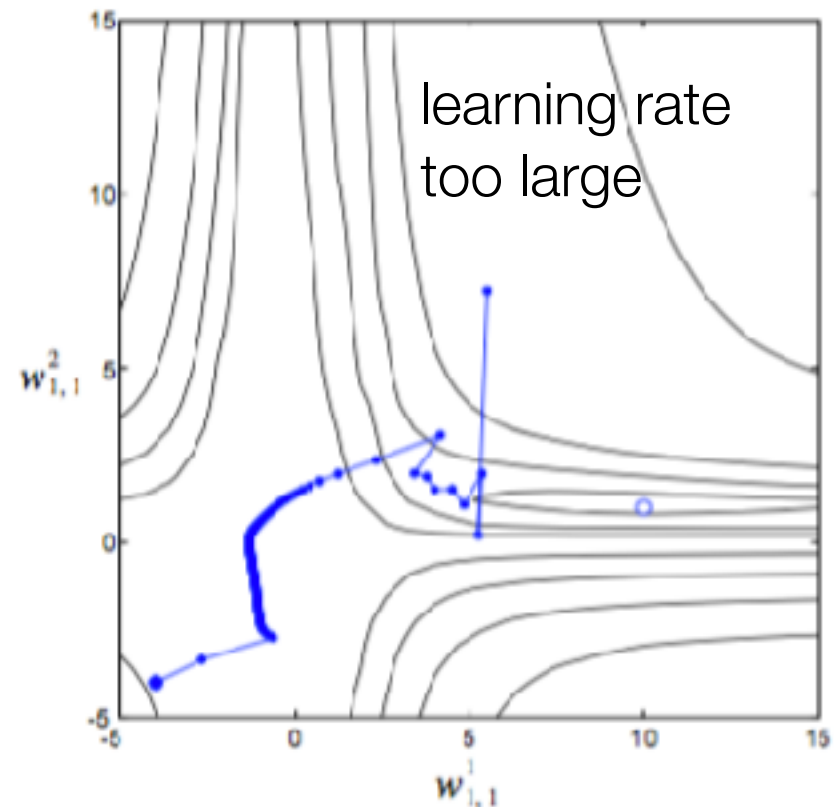
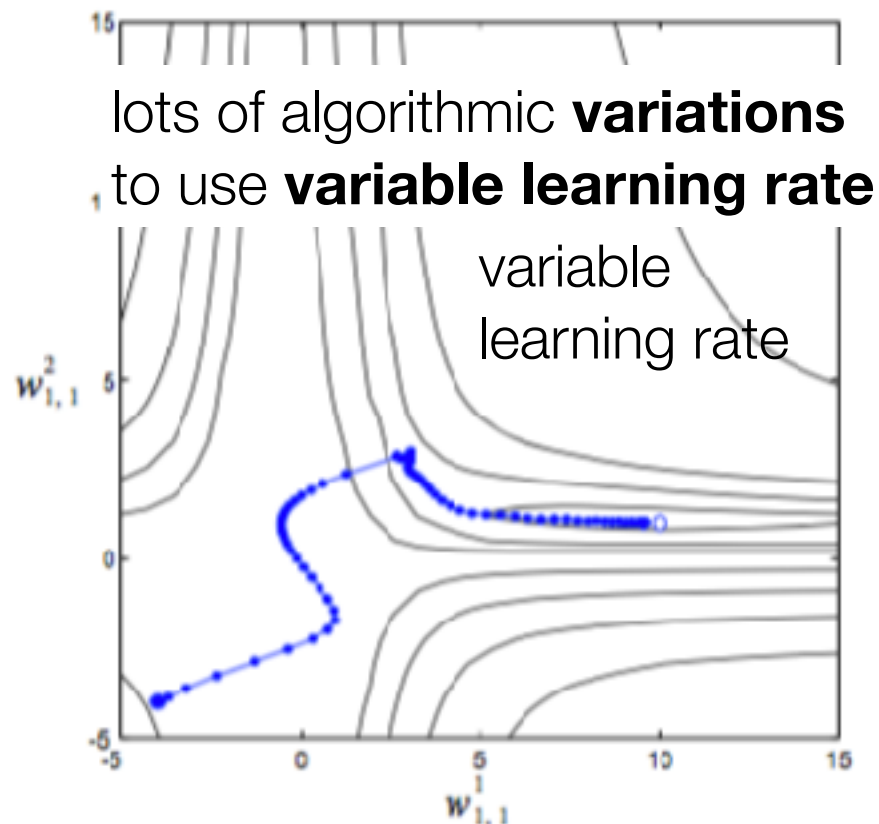
- Space is no longer convex
  - **One solution:**
    - start with large step size
    - “cool down” by decreasing step size for higher iterations



# Another Adaptive Strategy

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \eta^k \cdot \rho_k$$

- Space is no longer convex
  - **another solution:**
    - start with arbitrary step size
    - only decrease when successive iterations do not decrease cost



## 07. MLP Neural Networks.ipynb


### **comparison:**

mini-batch

momentum

adaptive learning


L-BFGS



# Fig. 2: NumPy is the base of the scientific Python ecosystem.


From: Array programming with NumPy

## Adaptive Optimization



# Be adaptive based on Gradient Magnitude?

- Decelerate down regions that are steep
- Accelerate on plateaus



Use **inverse** of  
**magnitude** of **gradient**!

Also **accumulate inverse** to be robust to  
**abrupt changes** in **steepness**...



# Common Adaptive Strategies

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Momentum

$$\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$$

- Nesterov's Accelerated Gradient

$$\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1}))}_{\text{step twice}} + \alpha \nabla J(\mathbf{W}_{k-1})$$

- AdaGrad

$$\rho_k = \frac{\eta}{\sqrt{G_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

where

$$G_k = G_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

all operations are per element

- RMSProp

$$\rho_k = \frac{\eta}{\sqrt{V_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

$$G_k = \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

$$V_k = \gamma \cdot V_{k-1} + (1 - \gamma) \cdot G_k$$

all operations are per element

- AdaDelta

$$\rho_k = \frac{\sqrt{M_k + \epsilon}}{\sqrt{V_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

$$M_k = \gamma \cdot M_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$

all operations are per element

- AdaM

$G$  updates with decaying momentum of  $J$  and  $J^2$

- NAdaM

same as Adam, but with nesterov's acceleration

**None** of these are “**one-size-fits-all**” because the space of neural network **optimization varies** by problem, ADAM is **popular** but **not a panacea**

# Adaptive Momentum

All operations are element wise:

$$\beta_1 = 0.9, \beta_2 = 0.999, \eta = 0.001, \epsilon = 10^{-8}$$

$$k = 0, \mathbf{M}_0 = \mathbf{0}, \mathbf{V}_0 = \mathbf{0}$$

Published as a conference paper at ICLR 2015

ADAM: A METHOD FOR STOCHASTIC OPTIMIZATION

Diederik P. Kingma\*  
University of Amsterdam, OpenAI

Jimmy Lei Ba\*  
University of Toronto

**For each epoch:**

update epoch  $k \leftarrow k + 1$

get gradient  $\mathbf{G}_k \leftarrow \nabla J(\mathbf{W}_k)$


accumulated gradient  $\mathbf{M}_k \leftarrow \beta_1 \cdot \mathbf{M}_{k-1} + (1 - \beta_1) \cdot \mathbf{G}_k$

accumulated squared gradient  $\mathbf{V}_k \leftarrow \beta_2 \cdot \mathbf{V}_{k-1} + (1 - \beta_2) \cdot \mathbf{G}_k \odot \mathbf{G}_k$

boost moments magnitudes  
(notice  $k$  in exponent)  $\hat{\mathbf{M}}_k \leftarrow \frac{\mathbf{M}_k}{(1 - [\beta_1]^k)} \quad \hat{\mathbf{V}}_k \leftarrow \frac{\mathbf{V}_k}{(1 - [\beta_2]^k)}$

update gradient, normalized  
by second moment  
similar to AdaDelta  $\mathbf{W}_k \leftarrow \mathbf{W}_{k-1} - \eta \cdot \frac{\hat{\mathbf{M}}_k}{\sqrt{\hat{\mathbf{V}}_k + \epsilon}}$

# Visualization of Optimization



<https://ruder.io/optimizing-gradient-descent/>