

# Lecture Notes for **Machine Learning in Python**

Professor Eric Larson  
**Neural Network Optimization and Activation**

# Class Logistics and Agenda

- Agenda:
  - More optimization and architectures
  - Programming Examples

# Last Time

# Problems with Advanced Architectures

- Numerous weights to find gradient update
  - minimize number of instances
  - **solution:** mini-batch
- **new problem:** mini-batch gradient can be erratic
  - **solution:** momentum
    - use previous update in current update

# Common Adaptive Strategies

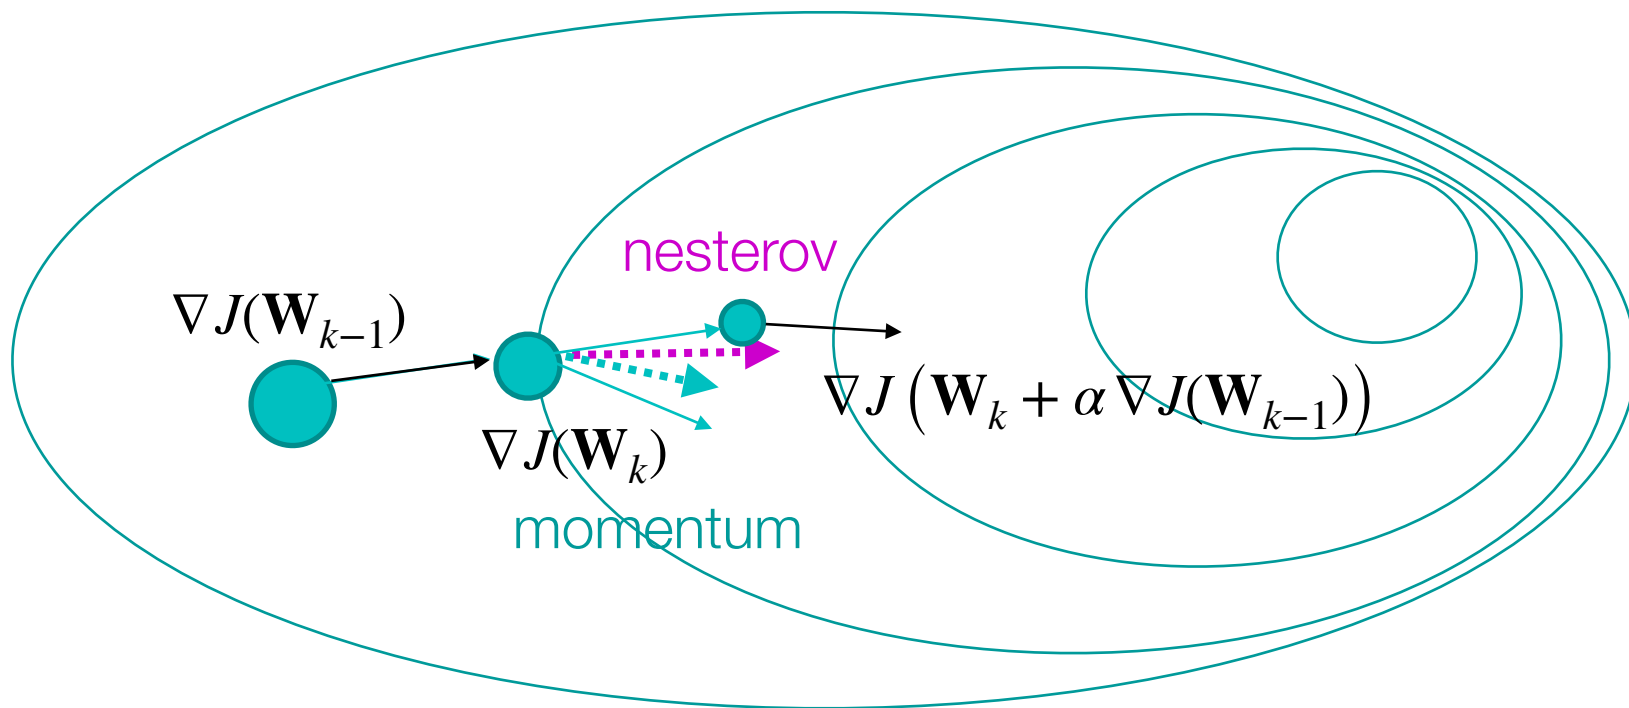
$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Momentum

$$\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$$

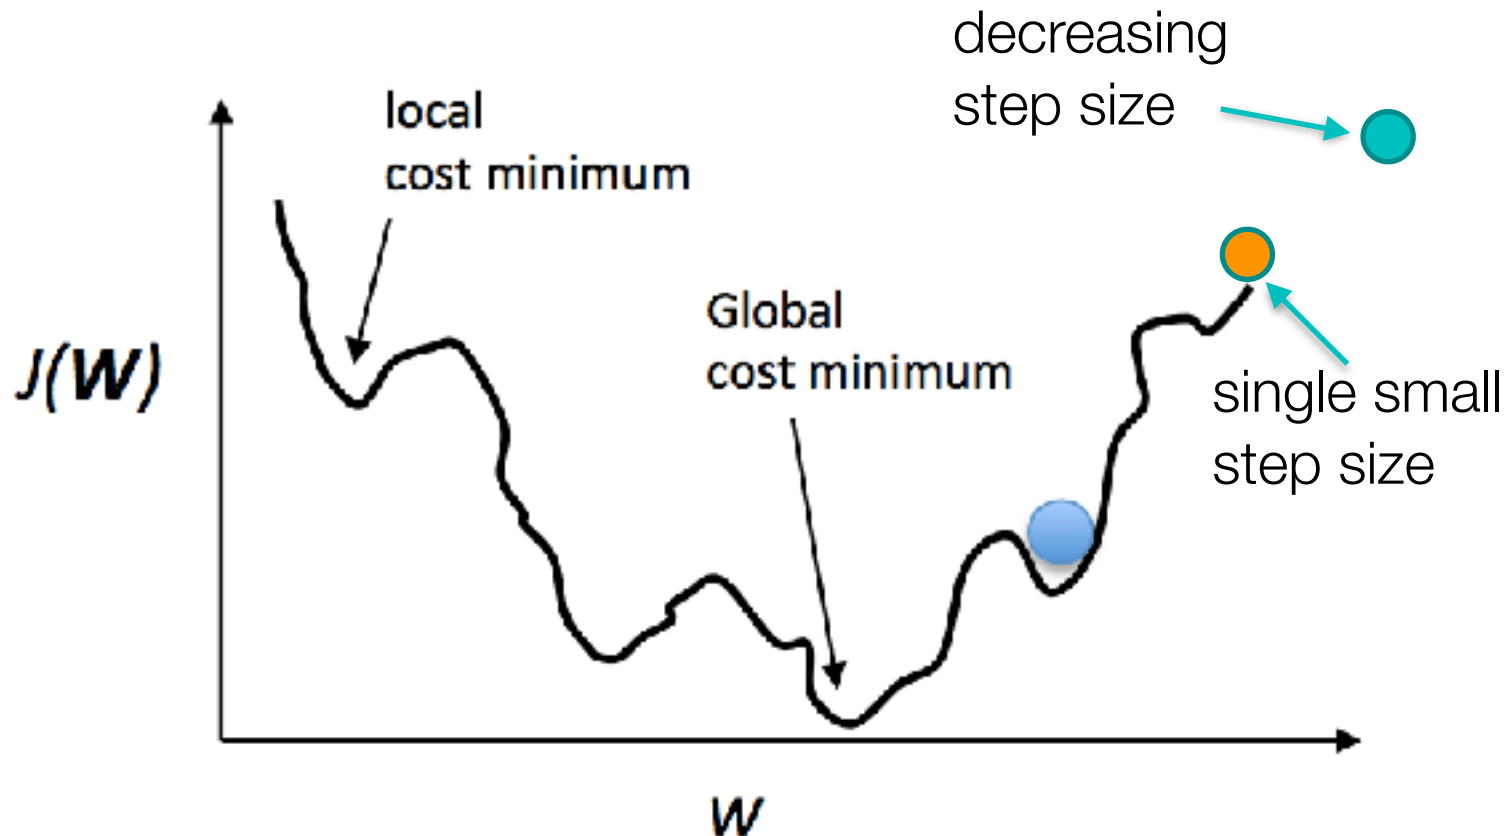
- Nesterov's Accelerated Gradient

$$\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1}))}_{\text{step twice}} + \alpha \nabla J(\mathbf{W}_{k-1})$$



# Adaptive Strategy: Cooling

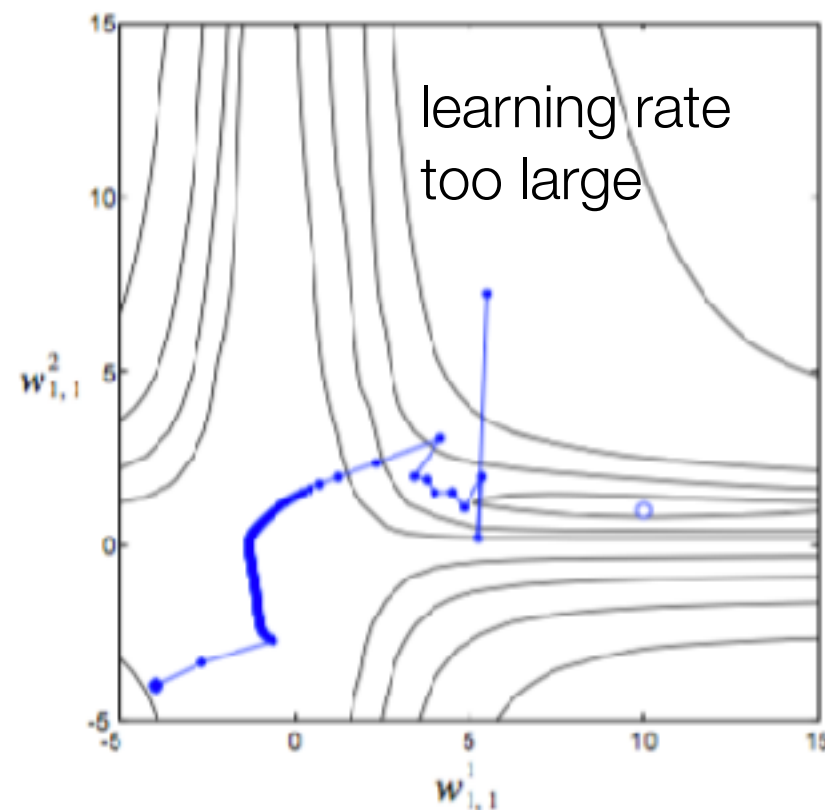
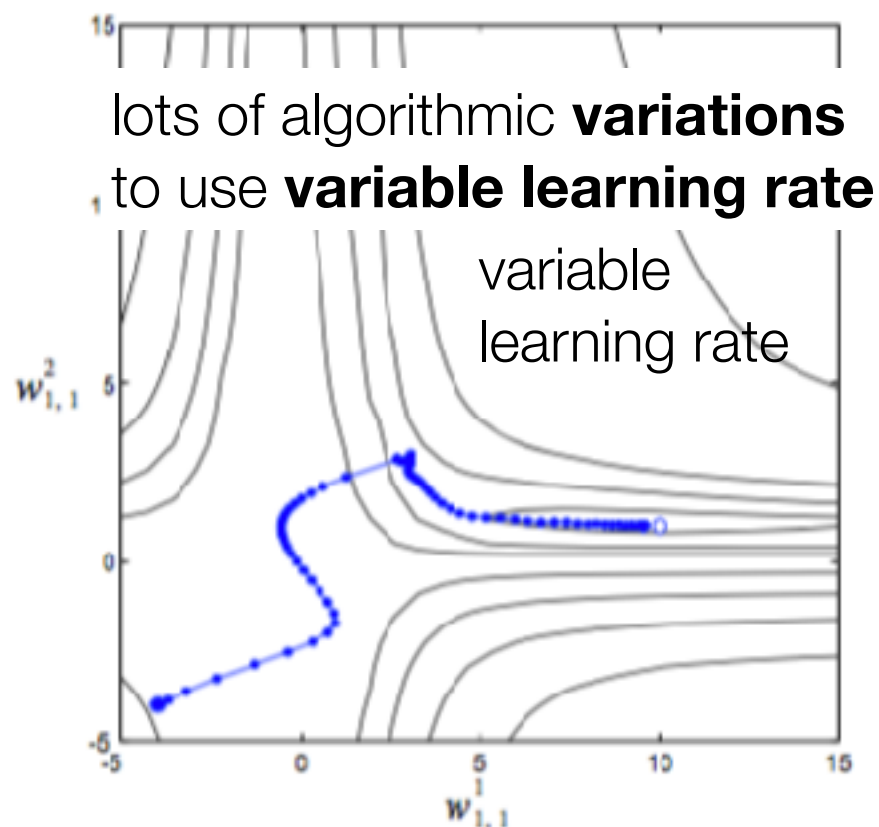
- Space is no longer convex
  - **One solution:**
    - start with large step size
    - “cool down” by decreasing step size for higher iterations



# Adaptive Strategies

- Space is no longer convex
  - **another solution:**
    - start with arbitrary step size
    - only decrease when successive iterations do not decrease cost

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \frac{\eta}{1 + \epsilon \cdot k} \cdot \rho_k$$



## 07. MLP Neural Networks.ipynb

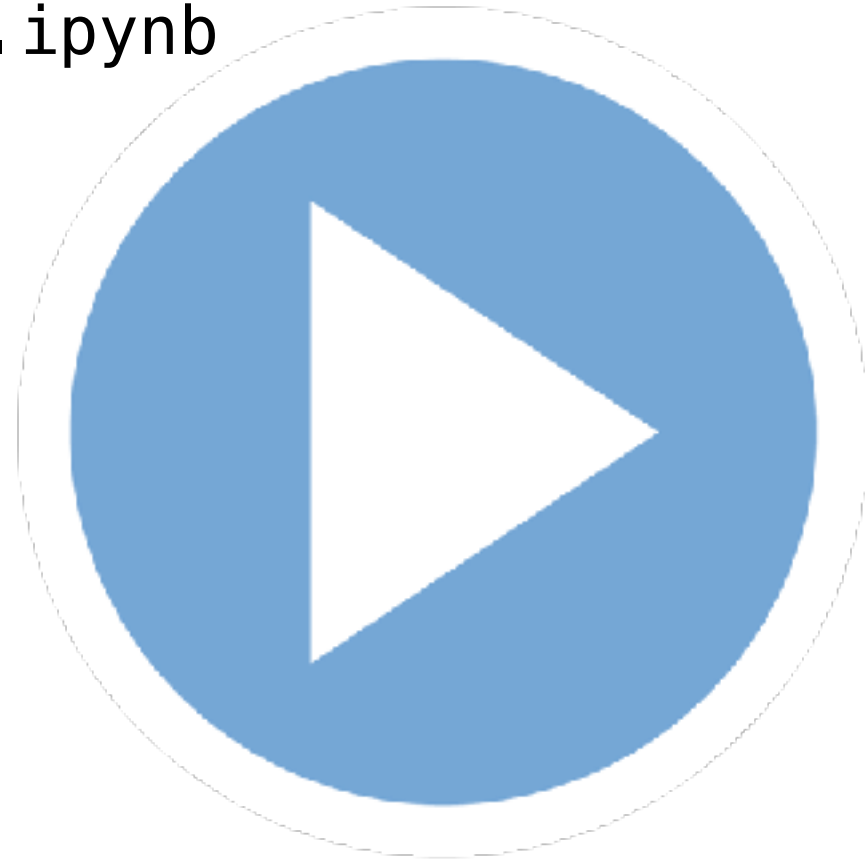
### **comparison:**

mini-batch

momentum

adaptive learning

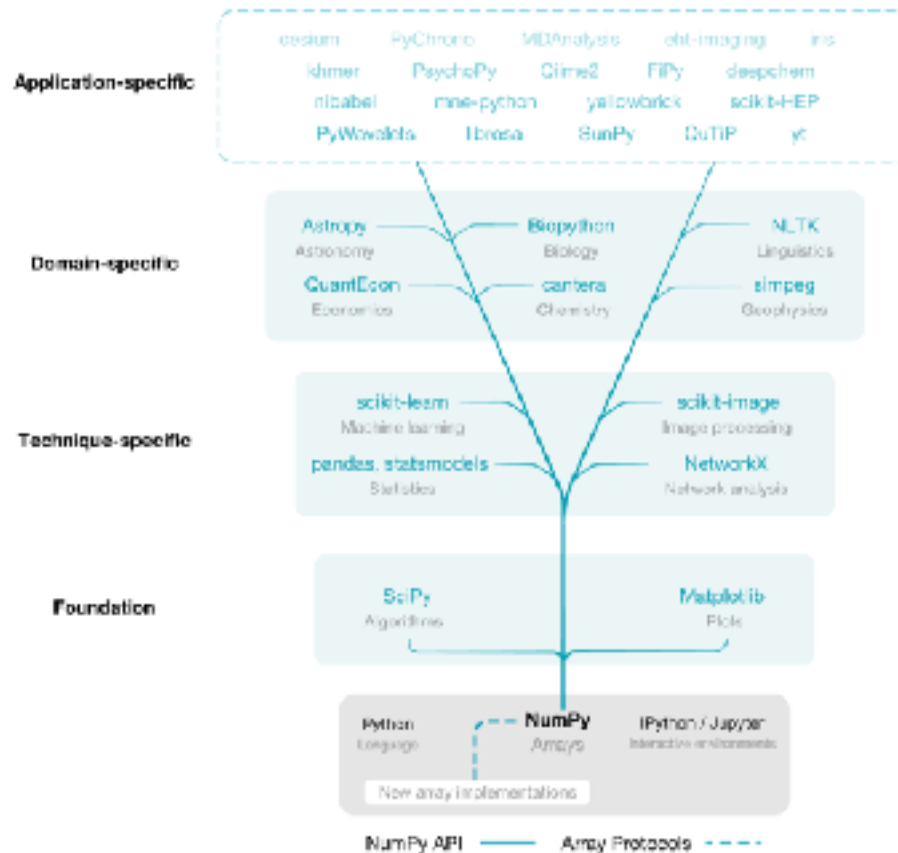
L-BFGS





# Fig. 2: NumPy is the base of the scientific Python ecosystem.

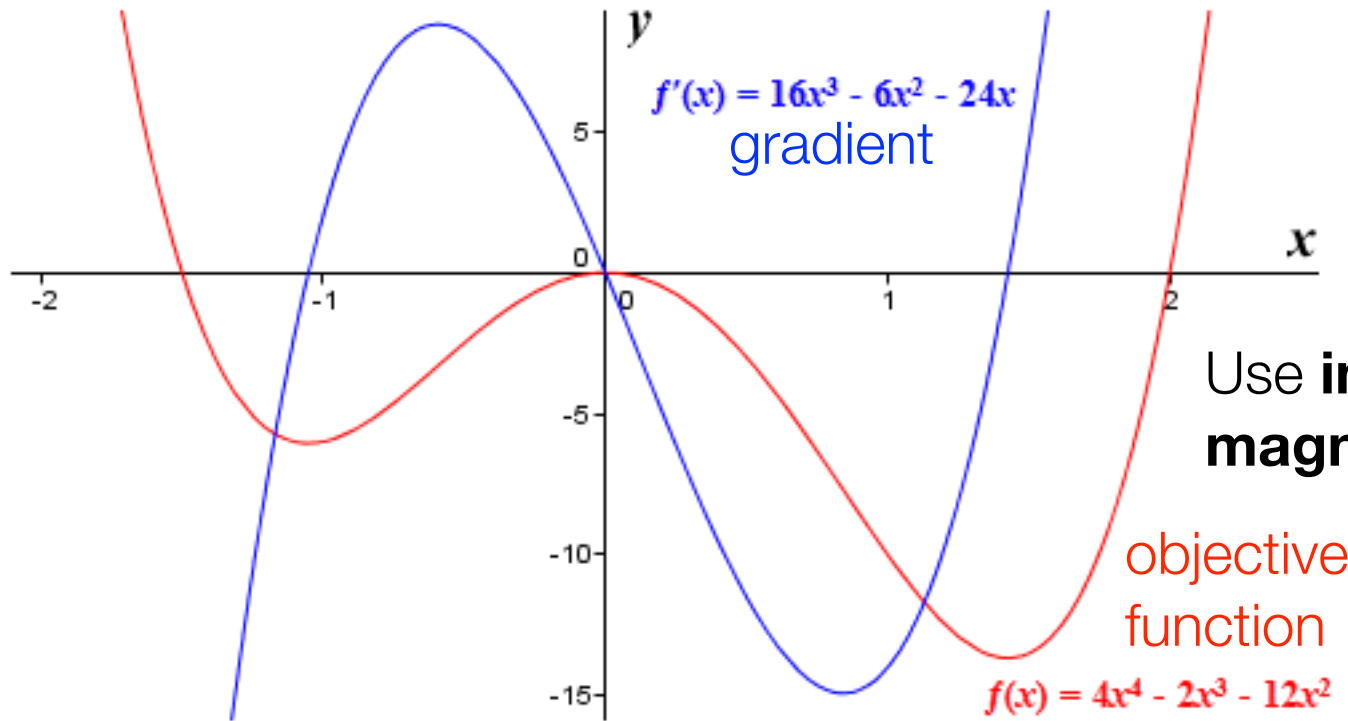
From: Array programming with NumPy



# Adaptive Optimization

# Be adaptive based on Gradient Magnitude?

- Decelerate down regions that are steep
- Accelerate on plateaus



Use **inverse** of  
**magnitude** of **gradient**!

Also **accumulate inverse** to be robust to  
**abrupt changes** in **steepness**...

# Common Adaptive Strategies

$$\mathbf{W}_{k+1} = \mathbf{W}_k - \rho_k$$

- Momentum

$$\rho_k = \alpha \nabla J(\mathbf{W}_k) + \beta \nabla J(\mathbf{W}_{k-1})$$

- Nesterov's Accelerated Gradient

$$\rho_k = \underbrace{\beta \nabla J(\mathbf{W}_k + \alpha \nabla J(\mathbf{W}_{k-1}))}_{\text{step twice}} + \alpha \nabla J(\mathbf{W}_{k-1})$$

- AdaGrad

$$\rho_k = \frac{\eta}{\sqrt{G_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \text{where} \quad G_k = G_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

all operations are per element

- RMSProp

$$\rho_k = \frac{\eta}{\sqrt{V_k + \epsilon}} \odot \nabla J(\mathbf{W}_k) \quad \begin{aligned} G_k &= \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k) \\ V_k &= \gamma \cdot V_{k-1} + (1 - \gamma) \cdot G_k \end{aligned}$$

all operations are per element

- AdaDelta

$$\rho_k = \eta \frac{M_k}{\sqrt{V_k + \epsilon}} \quad M_k = \gamma \cdot M_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$

all operations are per element

- AdaM

$G$  updates with decaying momentum of  $J$  and  $J^2$

- NAdaM

same as Adam, but with nesterov's acceleration

**None** of these are “**one-size-fits-all**” because the space of neural network **optimization varies** by problem, ADAM is **popular** but **not a panacea**

# Adaptive Momentum

All operations are element wise:

$$\beta_1 = 0.9, \beta_2 = 0.999, \eta = 0.001, \epsilon = 10^{-8}$$

$$k = 0, \mathbf{M}_0 = \mathbf{0}, \mathbf{V}_0 = \mathbf{0}$$

Published as a conference paper at ICLR 2015

ADAM: A METHOD FOR STOCHASTIC OPTIMIZATION

Diederik P. Kingma\*  
University of Amsterdam, OpenAI

Jimmy Lei Ba\*  
University of Toronto

**For each epoch:**

update epoch  $k \leftarrow k + 1$

get gradient  $\mathbf{G}_k \leftarrow \nabla J(\mathbf{W}_k)$

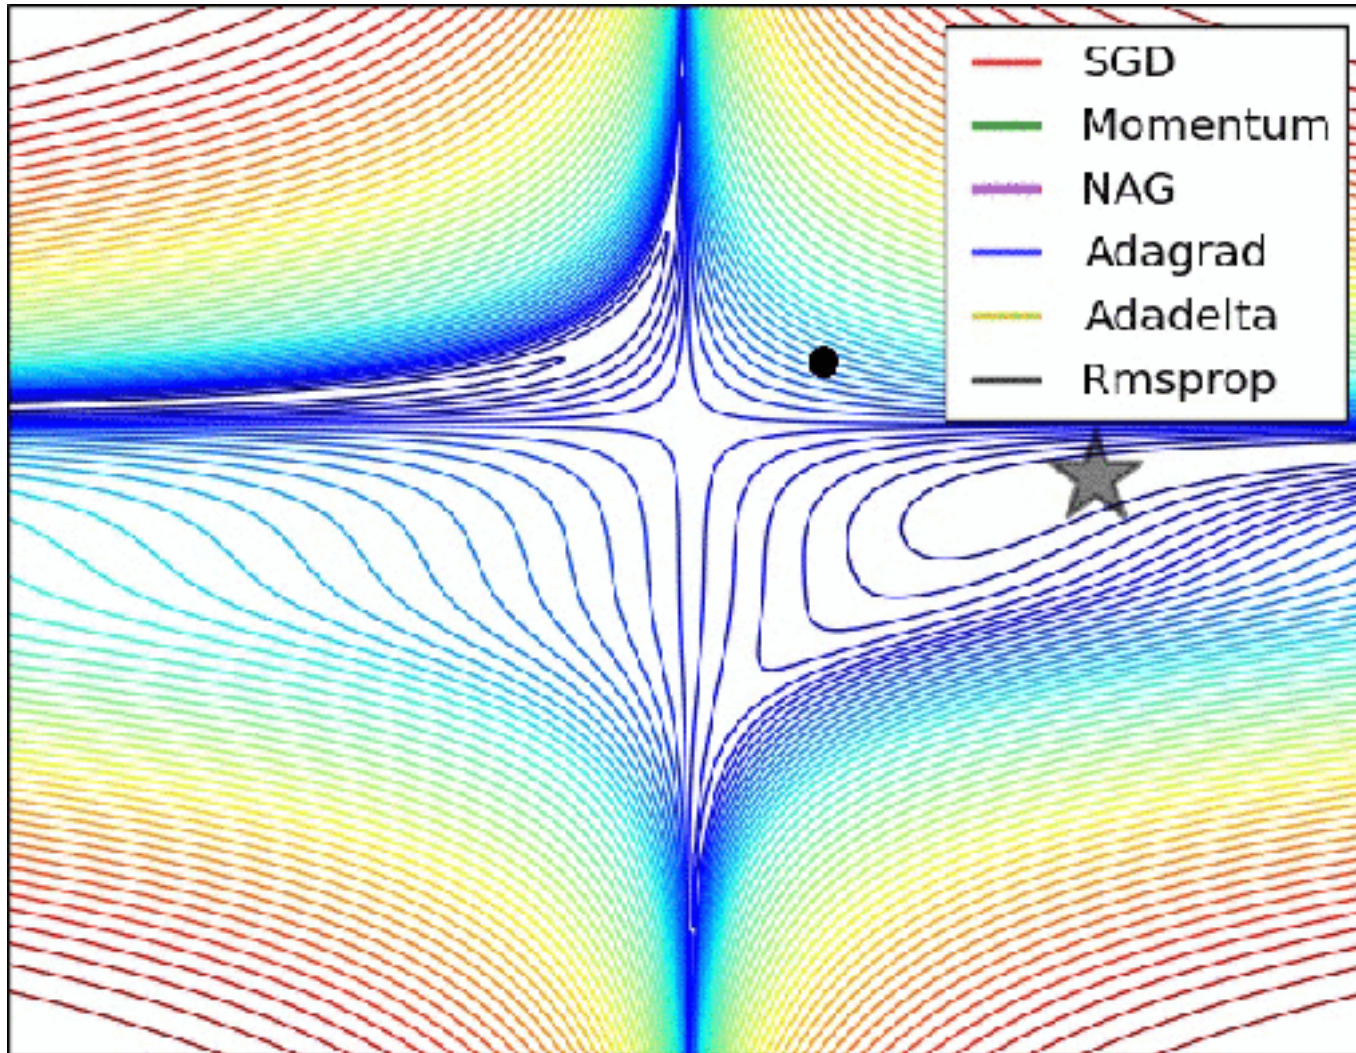
accumulated gradient  $\mathbf{M}_k \leftarrow \beta_1 \cdot \mathbf{M}_{k-1} + (1 - \beta_1) \cdot \mathbf{G}_k$

accumulated squared gradient  $\mathbf{V}_k \leftarrow \beta_2 \cdot \mathbf{V}_{k-1} + (1 - \beta_2) \cdot \mathbf{G}_k \odot \mathbf{G}_k$

boost moments magnitudes  
(notice  $k$  in exponent)  $\hat{\mathbf{M}}_k \leftarrow \frac{\mathbf{M}_k}{(1 - [\beta_1]^k)} \quad \hat{\mathbf{V}}_k \leftarrow \frac{\mathbf{V}_k}{(1 - [\beta_2]^k)}$

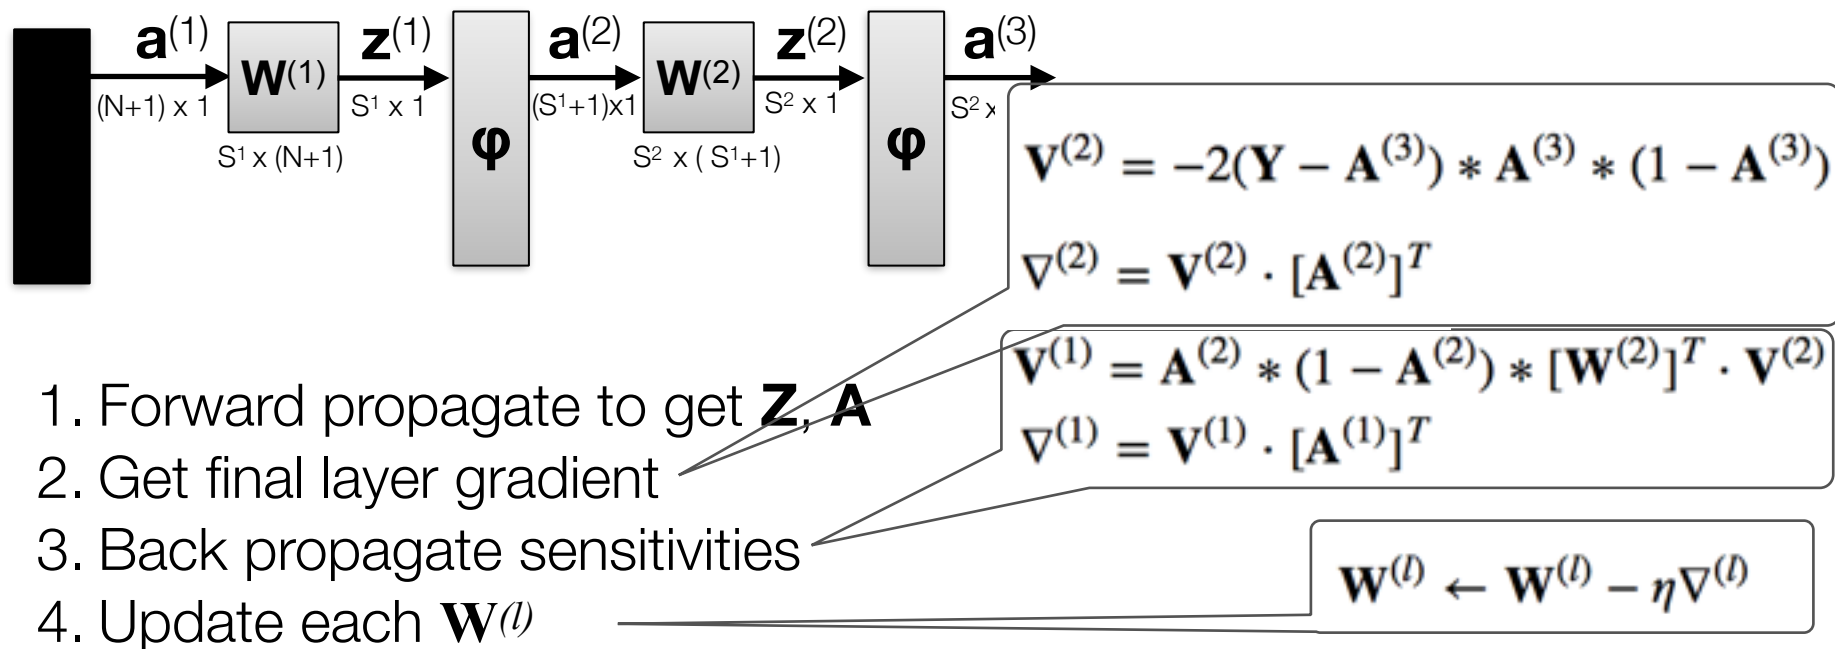
update gradient, normalized  
by second moment  
similar to AdaDelta  $\mathbf{W}_k \leftarrow \mathbf{W}_{k-1} - \eta \cdot \frac{\hat{\mathbf{M}}_k}{\sqrt{\hat{\mathbf{V}}_k + \epsilon}}$

# Visualization of Optimization



<https://ruder.io/optimizing-gradient-descent/>

# Changing the Objective Function



## • Self Test:

**True or False:** If we change the cost function,  $J(\mathbf{W})$ , we only need to update the final layer sensitivity calculation,  $\mathbf{V}^{(2)}$ , of the back propagation steps. The remainder of the algorithm is unchanged.

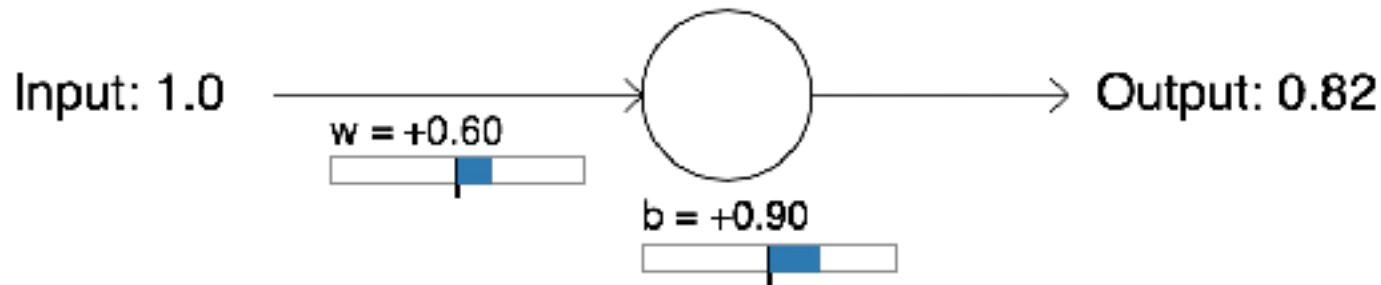
- A. True
- B. False

# Practical Implementation of Architectures

- MSE

$$J(\mathbf{W}) = \sum_k^M (\mathbf{y}^{(k)} - \mathbf{a}^{(L)})^2$$

least squares objective,  
tends to slow training initially



Run

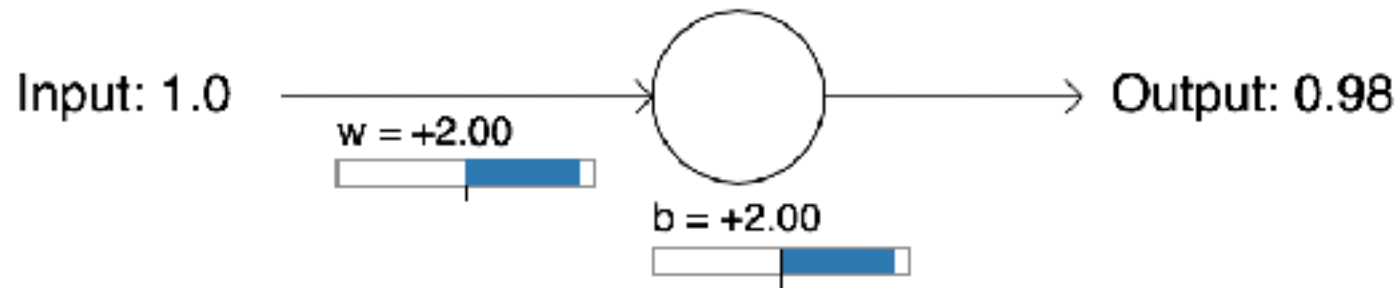


# Practical Implementation of Architectures

- MSE

$$J(\mathbf{W}) = \sum_k^M (\mathbf{y}^{(k)} - \mathbf{a}^{(L)})^2$$

least squares objective,  
tends to slow training initially



Run

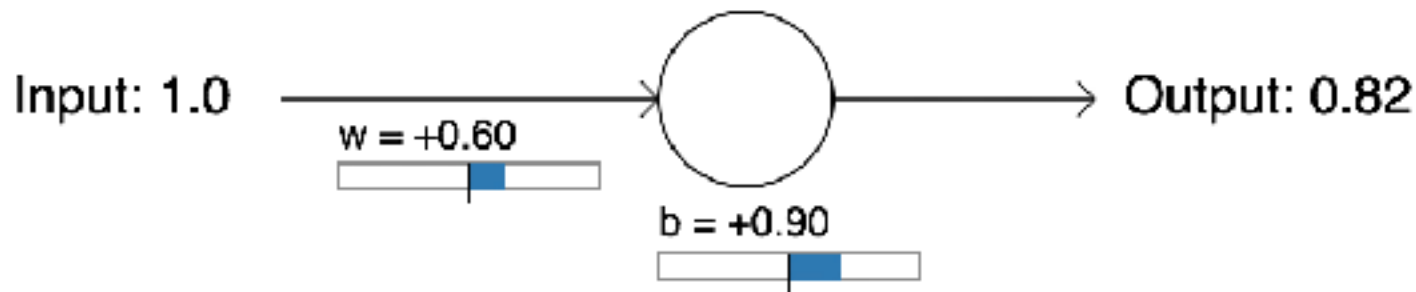


# Practical Implementation of Architectures

- Negative of MLE: **Binary Cross entropy**

$$J(\mathbf{W}) = - [\mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)})]$$

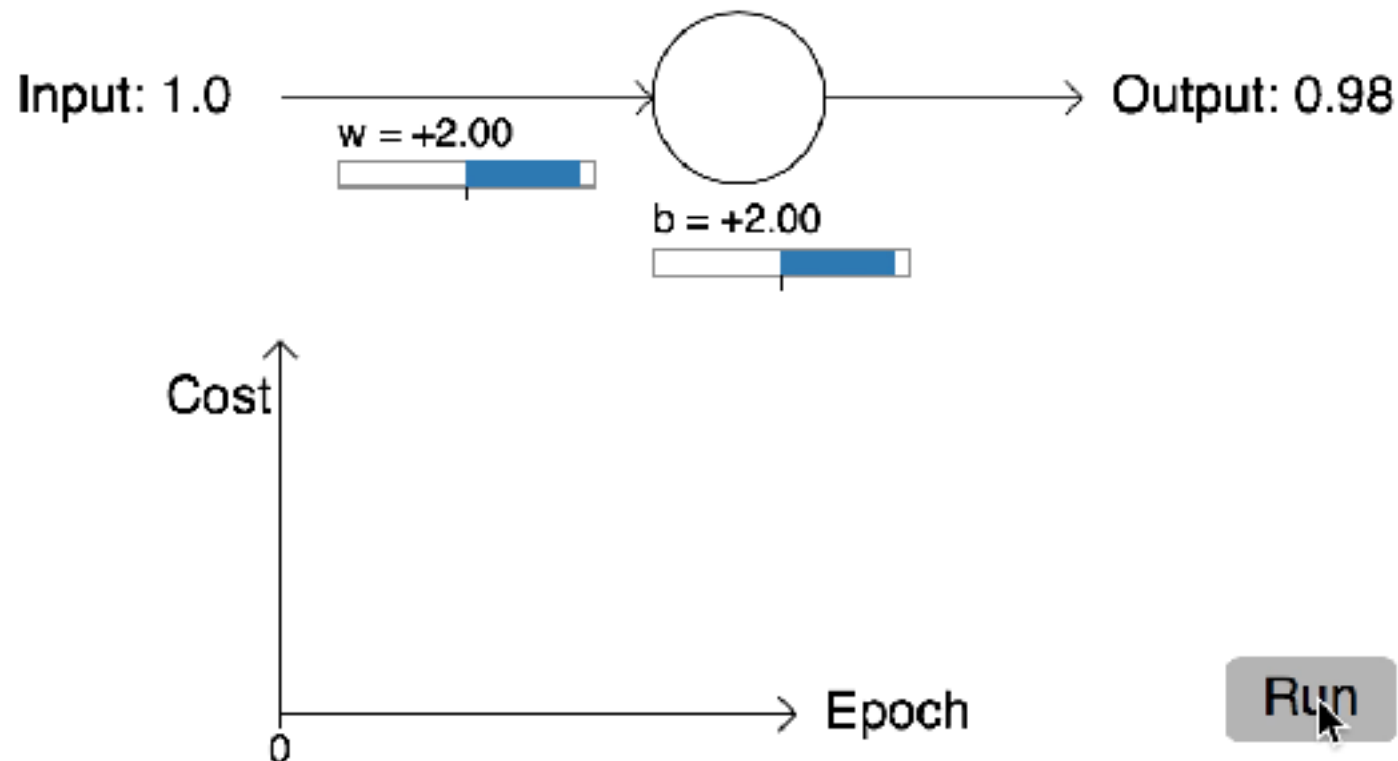
speeds up initial training



# Practical Implementation of Architectures

- Negative of MLE: **Binary Cross entropy**

$$J(\mathbf{W}) = - \left[ \mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)}) \right] \quad \text{speeds up initial training}$$



# Practical Implementation of Architectures

- Negative of MLE: **Binary Cross entropy**

$$J(\mathbf{W}) = - \left[ \mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)}) \right] \quad \text{speeds up initial training}$$

$$\left[ \frac{\partial J(\mathbf{W})}{\mathbf{z}^{(L)}} \right]^{(i)}$$

$$\mathbf{V}^{(2)} = -2(\mathbf{Y} - \mathbf{A}^{(3)}) * \mathbf{A}^{(3)} * (1 - \mathbf{A}^{(3)}) \quad \text{old update}$$

# Practical Implementation of Architectures

- Back to our old friend: **Cross entropy**

$$J(\mathbf{W}) = - \left[ \mathbf{y}^{(i)} \ln([\mathbf{a}^{(L+1)}]^{(i)}) + (1 - \mathbf{y}^{(i)}) \ln(1 - [\mathbf{a}^{(L+1)}]^{(i)}) \right]$$

speeds up  
initial training

$$\left[ \frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(L)}} \right]^{(i)} = ([\mathbf{a}^{(L+1)}]^{(i)} - \mathbf{y}^{(i)})$$

$$\left[ \frac{\partial J(\mathbf{W})}{\partial \mathbf{z}^{(2)}} \right]^{(i)} = ([\mathbf{a}^{(3)}]^{(i)} - \mathbf{y}^{(i)})$$

$$\mathbf{V}^{(2)} = \mathbf{A}^{(3)} - \mathbf{Y}$$

new update

```
# vectorized backpropagation
V2 = (A3 - Y_enc) # <- this is only line t
V1 = A2 * (1 - A2) * (W2.T @ V2)

grad2 = V2 @ A2.T
grad1 = V1[1:,:] @ A1.T
```

$$\mathbf{V}^{(2)} = -2(\mathbf{Y} - \mathbf{A}^{(3)}) * \mathbf{A}^{(3)} * (1 - \mathbf{A}^{(3)})$$

old update

bp-5

cross entropy

