

Lecture Notes for **Machine Learning in Python**

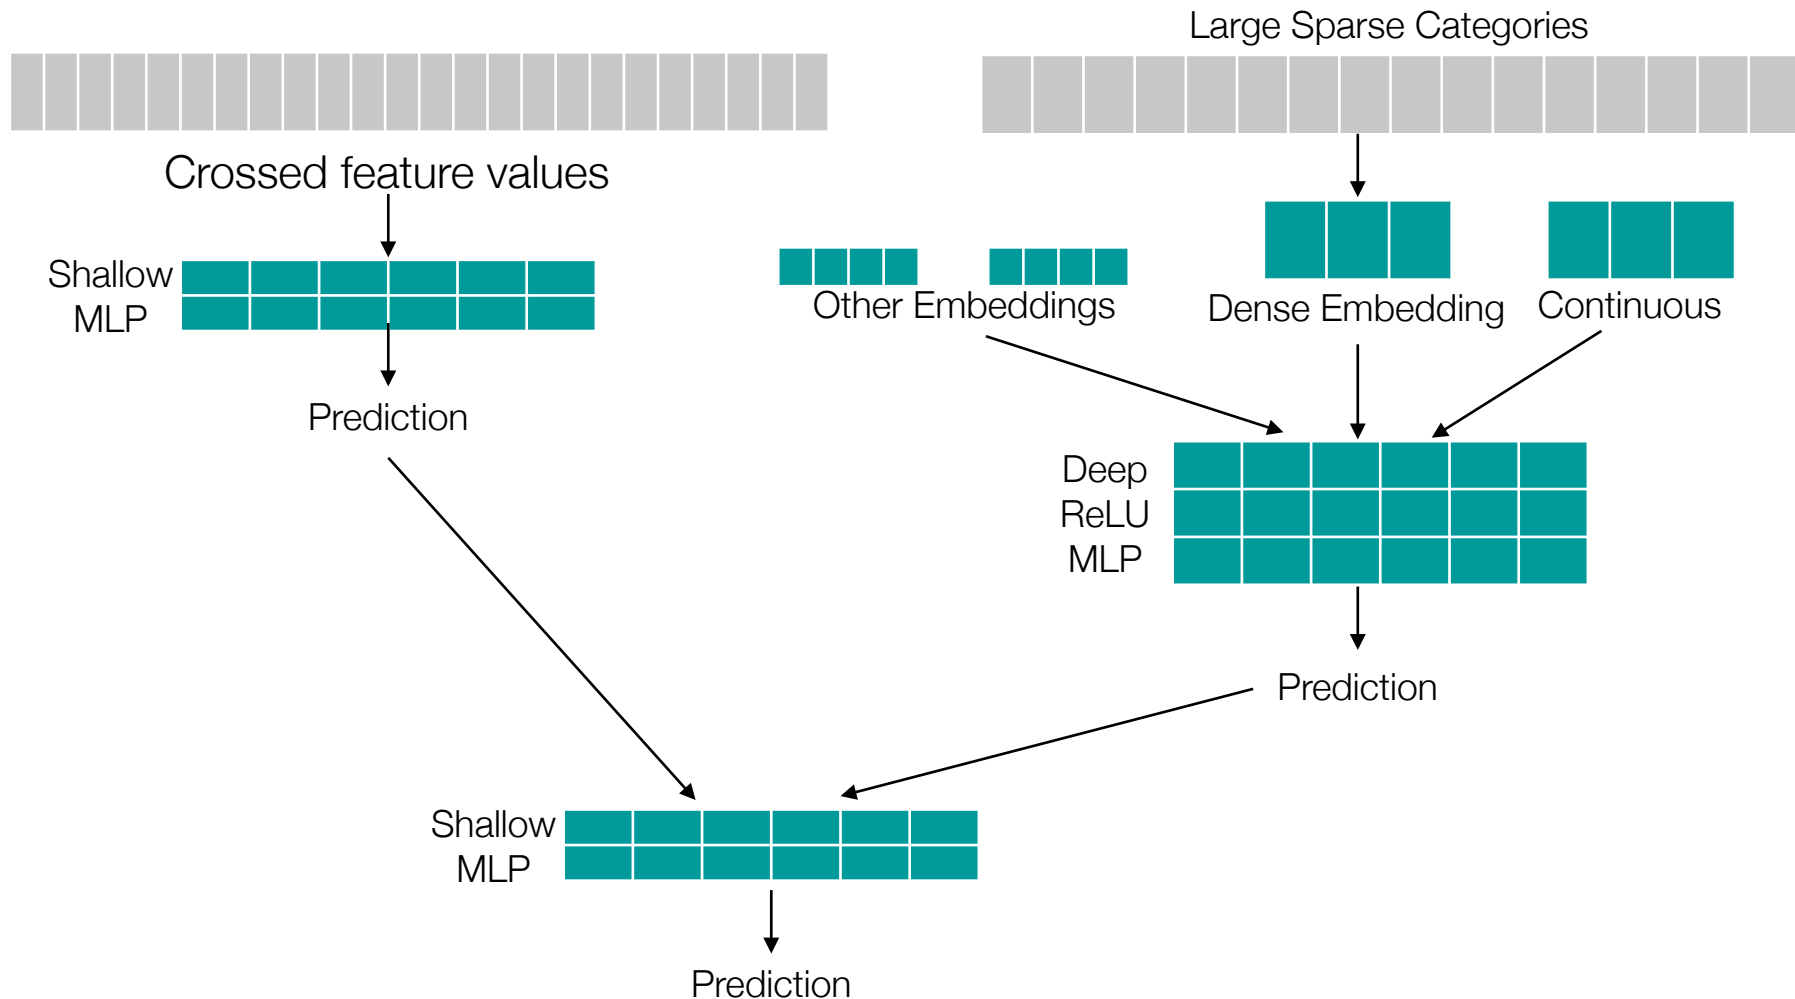
Professor Eric Larson
Basic Convolutional Neural Networks

Logistics and Agenda

- Logistics
 - Wide/Deep due soon!
 - Remember: Feel free to turn in late for partial credit.
- Agenda
 - Wide/Deep Finish Demo and Town Hall
 - Basic CNN architectures and Demo

Last Time:

- Deep refers to increasingly smaller hidden layers
- Embed into sparse representations via ReLU




10. Keras Wide and Deep.ipynb

The awful dataset:
Toy Census Data Example


Other tutorials:

https://www.tensorflow.org/tutorials/wide_and_deep



Town Hall, Wide and Deep Networks

When $p < 0.05$



Convolutional Neural Networks



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

Reminder: Convolution

$$\sum \left(\mathbf{I} \left[i \pm \frac{r}{2}, j \pm \frac{c}{2} \right] \odot \mathbf{k} \right) = \mathbf{O}[i, j] \quad \begin{array}{l} \text{output image} \\ \text{at pixel } i, j \end{array}$$

input image at $r \times c$ range of
pixels centered in i, j

kernel of size, $r \times c$
usually $r=c$

0	0	0	0	0	0	0	0	0
0	1	2	3	4	12	9	8	0
0	5	2	3	4	12	9	8	0
0	5	2	1	4	10	9	8	0
0	7	2	1	4	12	7	8	0
0	7	2	1	4	14	9	8	0
0	5	2	3	4	12	7	8	0
0	5	2	1	4	12	9	8	0
0	0	0	0	0	0	0	0	0

input image, \mathbf{I}

1	2	1
2	4	2
1	2	1

kernel
filter, \mathbf{k}
3x3

20	21	36
...
...
...
...
...
...
...

output image, \mathbf{O}

Reminder: Convolution

$$\begin{array}{|c|c|} \hline O_{11} & O_{12} \\ \hline O_{21} & O_{22} \\ \hline \end{array} = \text{Convolution} \left(\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array}, \begin{array}{|c|c|} \hline F_{11} & F_{12} \\ \hline F_{21} & F_{22} \\ \hline \end{array} \right)$$

Input **X**
Filter **F**

X ₁₁	X ₁₂	X ₁₃
X ₂₁	X ₂₂	X ₂₃
X ₃₁	X ₃₂	X ₃₃

Input **X**

⊗

F ₁₁	F ₁₂
F ₂₁	F ₂₂

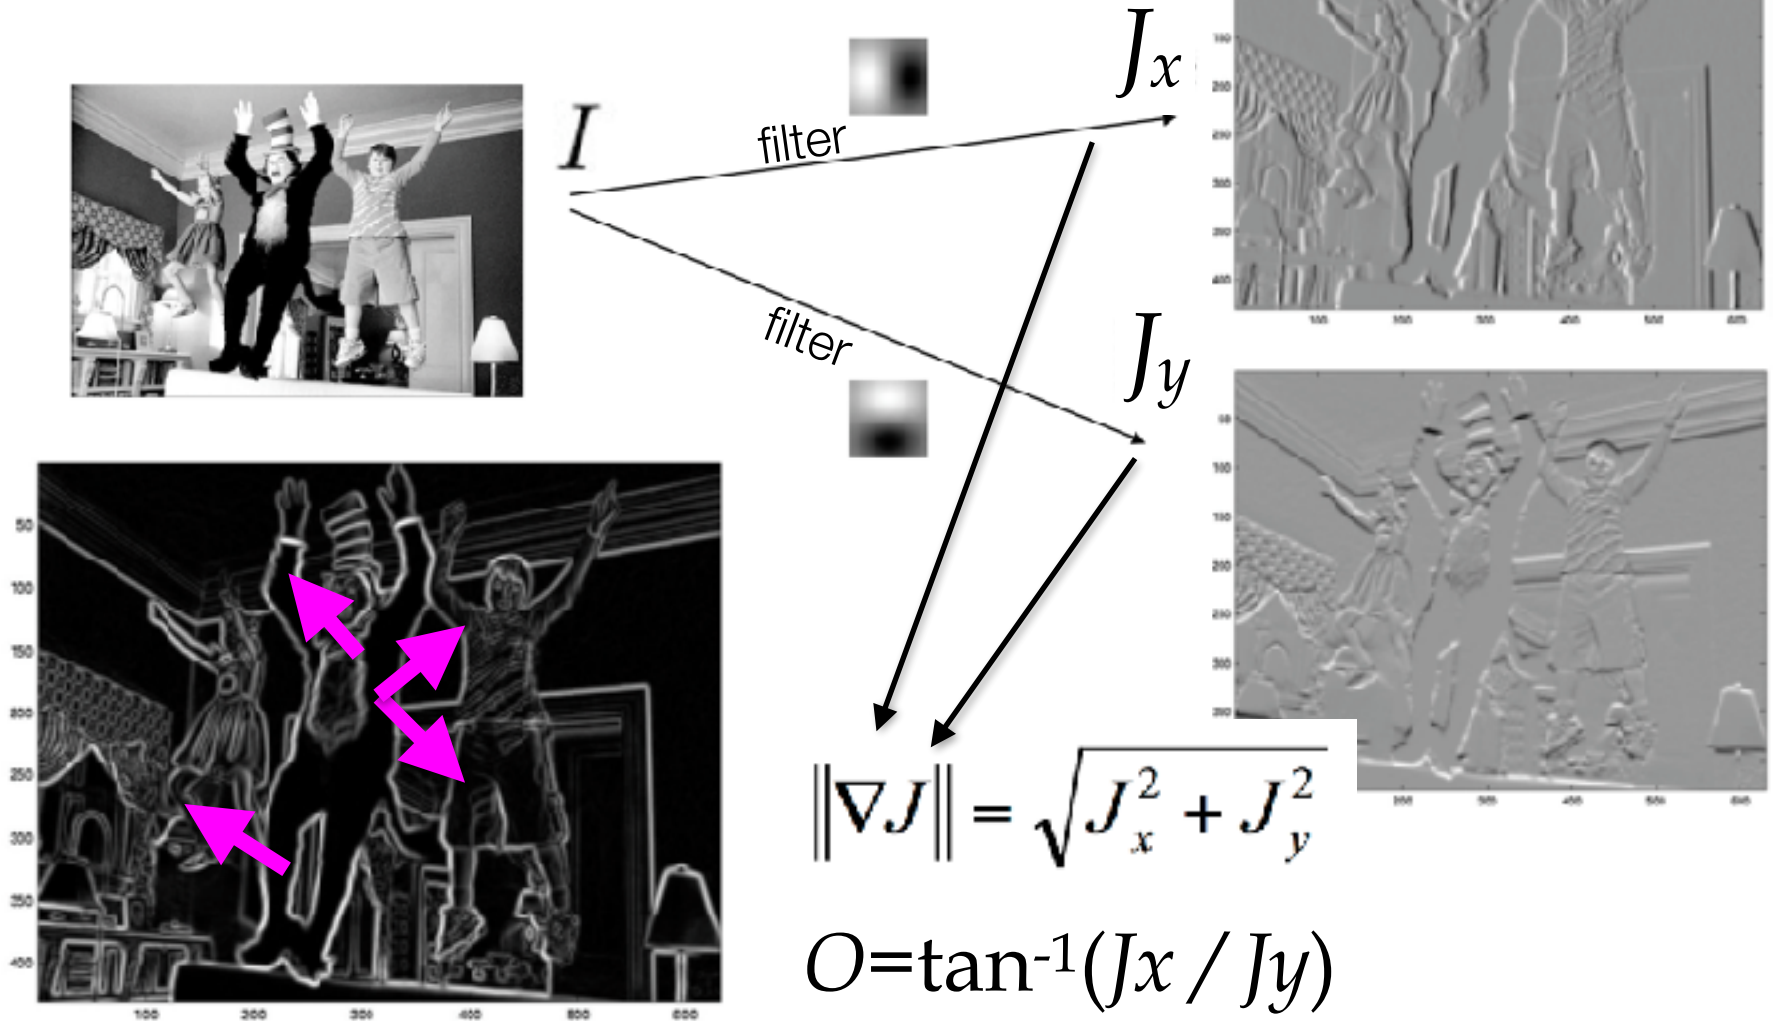
Filter **F**

X ₁₁ F ₁₁	X ₁₂ F ₁₂	X ₁₃
X ₂₁ F ₂₁	X ₂₂ F ₂₂	X ₂₃
X ₃₁	X ₃₂	X ₃₃

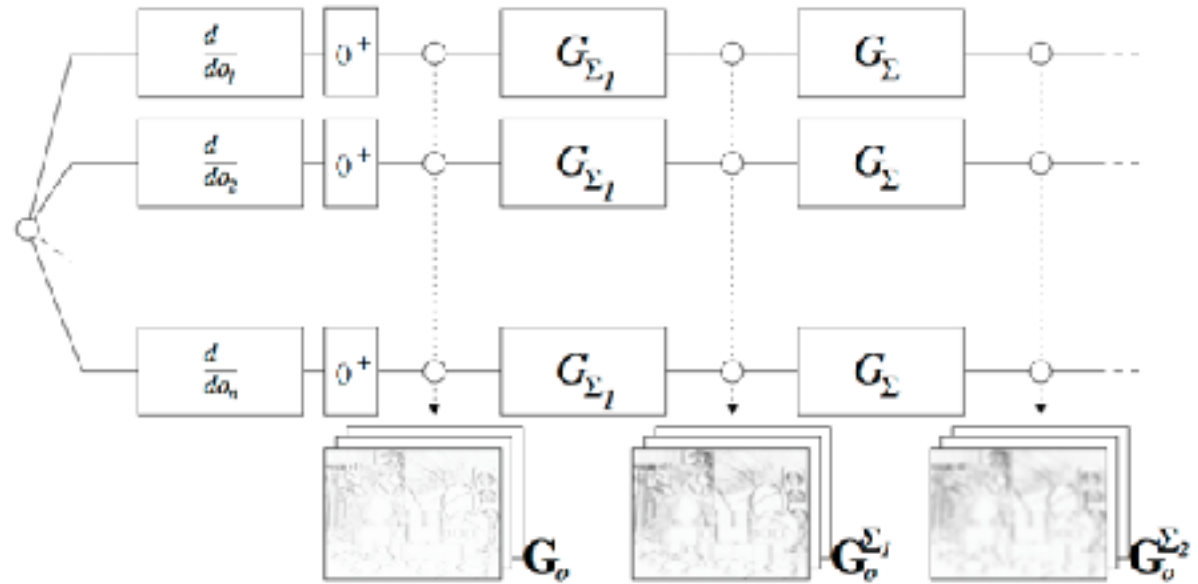
$$O_{11} = X_{11}F_{11} + X_{12}F_{12} + X_{21}F_{21} + X_{22}F_{22}$$

What we did before

- the gradient (2D derivative)



What we did before




take normalized histogram at point u, v

$$\tilde{\mathbf{h}}_{\Sigma}(u, v) = \left\| \left[\mathbf{G}_1^{\Sigma}(u, v), \dots, \mathbf{G}_H^{\Sigma}(u, v) \right]^{\top} \right\|$$

$$\mathcal{D}(u_0, v_0) = \begin{bmatrix} \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(u_0, v_0), \\ \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_1(u_0, v_0, R_1)), \dots, \tilde{\mathbf{h}}_{\Sigma_1}^{\top}(\mathbf{l}_T(u_0, v_0, R_1)), \\ \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_1(u_0, v_0, R_2)), \dots, \tilde{\mathbf{h}}_{\Sigma_2}^{\top}(\mathbf{l}_T(u_0, v_0, R_2)), \end{bmatrix}$$

Tola et al. "Daisy: An efficient dense descriptor applied to wide-baseline stereo." Pattern Analysis and Machine Intelligence, IEEE Transactions

Anatomy of a convolution




Blue Tensors: Outputs of Each Layer

Learned Params: Weights in Each Filter and Fully Connected Layer

CNN Overview

- First layer(s):
 - convolution
 - nonlinearity
 - pooling
 - Each pooling layer *can* make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons




CNN Overview: Self Test

- First layer(s):
 - convolution
 - nonlinearity
 - pooling
 - Each pooling layer *can* make the input image “smaller”
 - allows for “Information Distillation”
 - less dependence on exact pixels
- Final layers are densely connected
 - typically multi-layer perceptrons
- Where are unstable gradients **most** problematic?
 - (A) During Convolution Layer(s) updates
 - (B) During Fully Connected Layer(s) updates
 - (C) Both A and B
 - (D) They are not a problem


CNN Filtering

- Why perform lots of filtering?
 - “recall” gabor filtering?



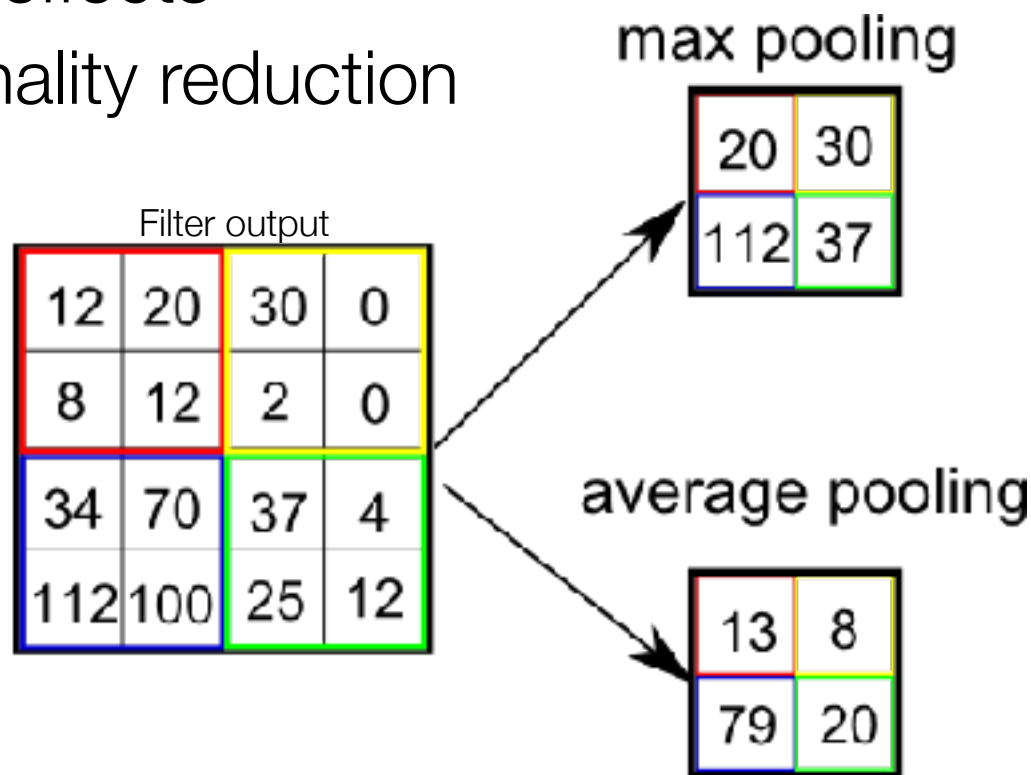


V1	Motion
V2	Stereo
V3	Color
V3a	Texture segregation
V3b	Segmentation, grouping
V4	Recognition
V7	Face recognition
MT	Attention
MST	Working memory/mental imagery

Convolution in a CNN

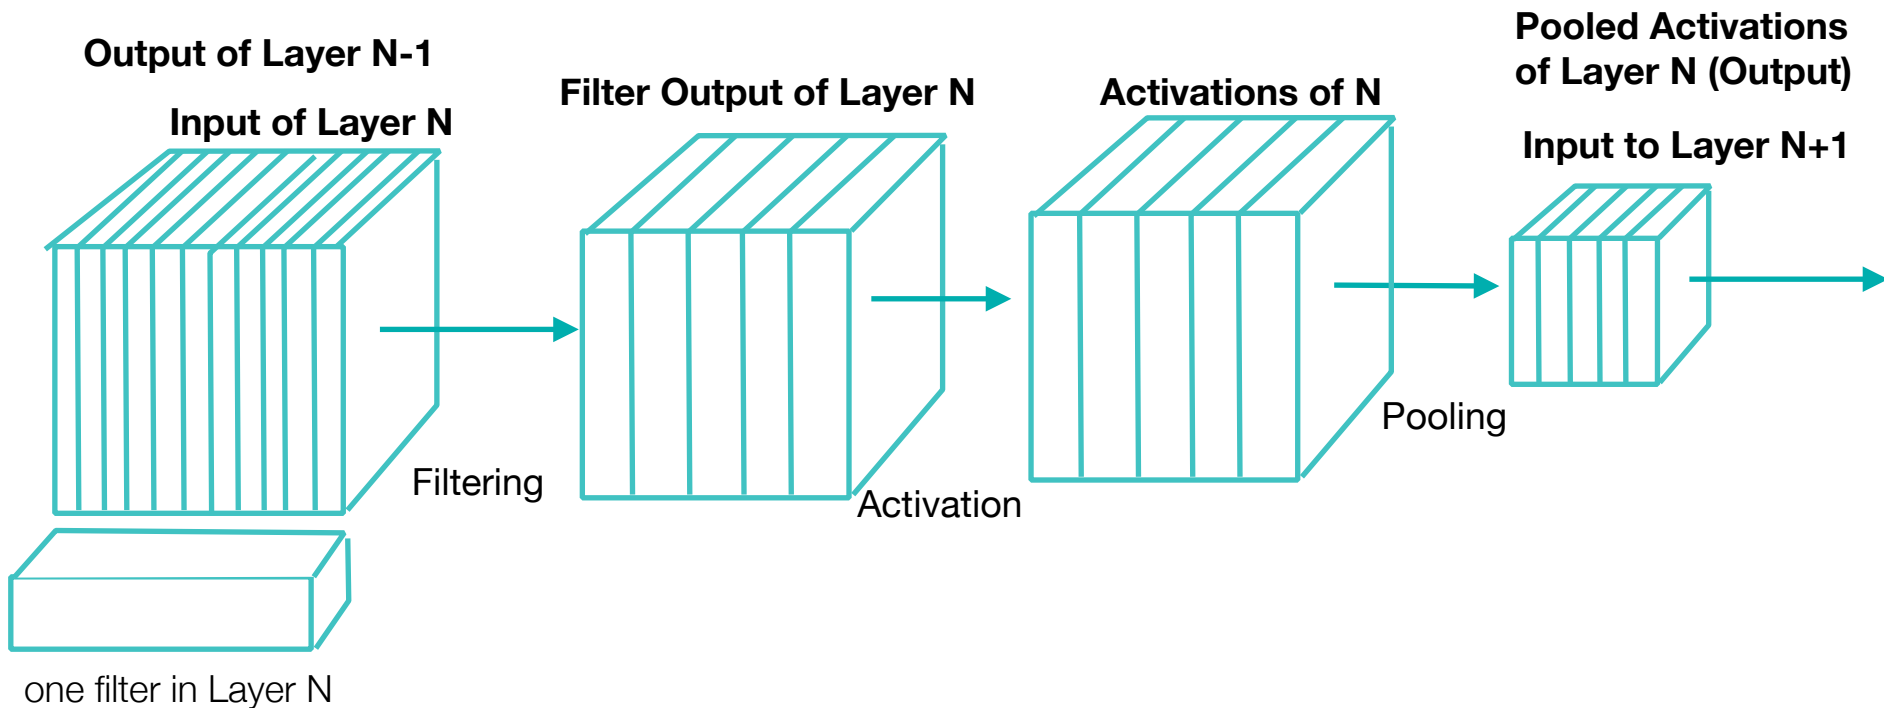


CNN Pooling

- Why perform pooling?
- Why max pooling?
 - reduce translation effects
 - **mostly**: dimensionality reduction



CNNs: Putting it together



Structure of Each Tensor: Channels x Rows x Columns